

Better Jobs for Mothers, More Sons at Birth^{*}

Nghiêm Q. Huỳnh[†]

Ngoc T. T. Nguyen[‡]

May 2026

Abstract

Women's labor-market gains are often expected to favor daughters. This paper shows why they can instead intensify male-biased sex selection when son preference remains strong and childcare is primarily maternal. We develop a household model in which maternal wage gains raise the time cost of childbearing, lower fertility, and increase sex selection, even as higher female returns narrow gender gaps in education. We test these predictions in Vietnam using industry-level tariff cuts from the 2001 US-Vietnam Bilateral Trade Agreement, assigning exposure separately to mothers and fathers. Maternal exposure increases male birth probability and reduces fertility. Paternal exposure raises fertility but does not affect sex selection. These contrasts distinguish the maternal time-cost mechanism from income and paternal-side explanations, suggesting that closing gender gaps in labor markets and education can widen gender gaps at birth when mothers bear most childcare responsibilities.

JEL Classification: F13, J13, J16, J22

Keywords: Sex ratio at birth, Sex selection, Fertility, Female labor supply, Trade liberalization

^{*}We are grateful to Mark Rosenzweig for extensive comments and guidance that substantially improved the paper, and to Brian McCaig for advice on the VHLSS data. We also thank S. Anukriti, Klaus Desmet, Tim Guinane, Arthur Lewbel, Kaivan Munshi, and Mu-Jeung Yang for helpful comments and suggestions. We thank seminar and conference participants at Boston College, Yale, the University of Oklahoma, Fulbright University Vietnam, the International School at Vietnam National University (Hanoi), the Chinese Economists Society North America Annual Conference at the University of Michigan, and NEUDC at Tufts University for useful feedback. Huỳnh gratefully acknowledges financial support from the Yale Council on Southeast Asian Studies. All errors are our own.

[†]Department of Economics, University of Oklahoma, Norman, OK 73019, USA (e-mail: nqh@ou.edu)

[‡]Disney Streaming, Los Angeles, CA, USA (e-mail: jade.ng.nguyen@gmail.com)

1. Introduction

Vietnam’s sex ratio at birth rose from 107 boys per 100 girls in 1999 to 112 by 2015. Over the same period, female wage employment nearly doubled as a share of total female employment, fertility remained low, and women’s university enrollment rose relative to men’s (Figure 1). Similar combinations of rising sex ratios and expanding labor-market opportunities for women appear in other countries with son preference. An influential view holds that women’s economic gains often improve child outcomes and can benefit daughters in particular (Jayachandran and Voena, 2026; Duflo, 2003; Qian, 2008). Why, then, can closing gender gaps in jobs and schooling coincide with a widening gender gap at birth?


The sex-selection literature points in several directions, and the evidence is mixed. Fertility decline can raise the sex ratio when families with son preference have fewer births over which to realize a desired son (Jayachandran, 2017). Fertility restrictions can intensify the same pressure, as evidence from China links stricter fertility limits under the one-child policy to more male-biased sex ratios (Li et al., 2011; Ebenstein, 2011). A different prediction arises when economic opportunities improve for women. If these gains raise the economic value of daughters, relative-returns arguments predict that sex selection should fall (Rosenzweig and Schultz, 1982; Qian, 2008). By contrast, income gains can worsen sex ratios when sex selection rises with household resources, as in China’s land reform (Almond et al., 2019).¹ These mechanisms imply different responses to rising female employment, but the evidence often comes from shocks assigned at the location or household-income level. Such variation cannot distinguish whether a response comes from mothers’ labor-market opportunities, fathers’ income, or broader local gains. Without that separation, fertility decline, relative returns, and income remain difficult to distinguish empirically.

This paper separates these channels by comparing labor-market gains that reach mothers and fathers differently. We show that maternal labor-market gains increase the male share at birth, while paternal labor-market gains do not. We explain the contrast through a household model in which mothers bear childcare time costs and parents have son preference. The fertility literature has long emphasized that higher female wages can reduce fertility because childcare remains disproportionately maternal (Doepke et al., 2023). In a setting with sex selection, this familiar mechanism has an additional implication. When maternal wages rise, children become more expensive in time, fertility falls, and parents with son preference have stronger incentives to concentrate remaining births on sons. When paternal wages rise, household resources increase without tightening the maternal time constraint. The model therefore provides a microfoundation for the fertility-decline mechanism emphasized by Jayachandran (2017), while distinguishing that channel from income and relative-returns explanations.

Our model also explains why rising investment in girls need not imply less sex selection. We

¹Marriage-market channels provide another mechanism in settings where relative wealth and marriage payments are central (Borker et al., 2026). We do not model this channel because it is not first-order in Vietnam, where dowry is uncommon.

Figure 1: Sex Ratio, Female Wage Employment, Fertility, and Education Gender Parity



Notes: Data from [Chao et al. \(2019\)](#) (panel a) and the World Bank World Development Indicators (panels b–d). Vietnam is the solid black line. Vertical dashed line marks 2001 (BTA). Horizontal dotted line in panel (a) marks the natural sex ratio (≈ 105). Panel (d) is the Gender Parity Index for university enrollment (female-to-male ratio). Log scale on all vertical axes.

combine the time-cost margin with the quantity-quality tradeoff ([Becker and Lewis, 1973](#)). Higher returns to daughters could in principle reduce sex selection by raising the value of daughters ([Qian, 2008](#)). But these higher returns also raise desired investment per daughter. Panel (d) of Figure 1 shows the relevant fact for Vietnam: the female-to-male university enrollment ratio moves toward parity over the same period in which the sex ratio at birth becomes more male-biased. In the model, the narrowing education gap makes children more expensive to rear, reinforcing the fertility decline already operating through the time-cost margin. Under son preference, the extra fertility decline produces more male-skewed births. Combining the two margins shows how maternal labor-market gains, narrowing education gaps, and rising male shares at birth can coexist in the same setting.

We test these predictions using the 2001 US-Vietnam Bilateral Trade Agreement (BTA), which reduced US tariffs on Vietnamese exports by different amounts across industries. These industry-specific export-demand shocks reached mothers and fathers differently depending on where each parent worked. We combine the 1999 Vietnam Population Census with the Vietnam Household

Living Standards Surveys (VHLSS, 2002–2010), which record each parent’s industry and employment close to the time of birth. We restrict attention to children aged 0 and 1, so the observed sex ratio closely tracks the sex ratio at birth and is less affected by differential mortality, migration, or later household composition changes. Observing both parents’ industries at the time of birth is the key advantage of these data. Prior studies of sex selection often rely on censuses or retrospective birth histories that recover longer fertility series but cannot assign contemporaneous industry to both parents for each birth. If parents switch industries or migrate after conception, later measures make it harder to distinguish maternal exposure, paternal exposure, and broader provincial exposure. The VHLSS gives a shorter but sharper window: it identifies parent-specific exposure close to birth, allowing us to separate shocks operating through the mother’s job, the father’s job, and the broader local economy.

The estimates match the model’s discriminating pattern. At the median industry tariff cut, maternal exposure raises the probability of a male birth by about 2 percentage points. The binary high-exposure design gives the same conclusion: mothers in high-exposure industries are more likely to have sons after the BTA. These mothers are also 15 percentage points more likely to hold a wage job, relative to a base of 25 percent, and work about 15 percent more hours per month. Annual birth probability falls by 0.6–0.7 percentage points in the reconstructed birth-hazard panel. Exposure through fathers’ industries or local labor markets has no detectable effect on sex selection. Paternal exposure does, however, raise annual birth probability by about 0.5 percentage points. This asymmetry is consistent with the model’s prediction that paternal income operates without tightening the maternal time constraint.

The paper contributes to three literatures. First, the model reconciles findings in the sex-selection literature that appear to point in opposite directions. In China, [Almond et al. \(2019\)](#) show that income gains increased sex selection, while [Qian \(2008\)](#) finds that higher female-specific returns reduced it. Evidence on the one-child policy also shows how fertility constraints can move sex ratios by limiting the quantity response ([Li et al., 2011](#); [Ebenstein, 2011](#)). The key distinction is whether fertility can adjust and whose income changes. When fertility is constrained, income shocks are more likely to operate through birth composition. When female-specific returns rise, greater investment per child can make daughters more valuable. In less constrained settings such as Vietnam, fertility can adjust, and the maternal time-cost channel makes the source of income decisive. Empirically, this contribution comes from separating maternal and paternal exposure. Existing studies commonly identify sex-selection effects from geographic or household-level variation ([Qian, 2008](#); [Almond et al., 2019](#); [Anukriti and Kumler, 2019](#); [Anukriti et al., 2022](#); [Chakraborty, 2015](#)). Such shocks assign the same exposure to all households in a place, while our design uses industry-level exposure that varies within the household.

Second, our framework links the fertility and human-capital literatures. These literatures often treat the quantity-quality tradeoff and sex selection as separate margins ([Rosenzweig and Wolpin, 1980](#); [Rosenzweig and Zhang, 2009](#); [Jayachandran and Kuziemko, 2011](#)). The closest related pa-

per is [Anukriti et al. \(2022\)](#), who show that ultrasound diffusion in India increased prenatal sex selection while improving postnatal outcomes for surviving girls. Their shock changes access to sex-selection technology and studies postnatal investments among surviving children. Our shock is a demand-side labor-market change that, combined with parent-specific industry data, allows us to test whether the source of labor-market gains affects fertility and sex composition at birth.

Third, the paper contributes to the Vietnam BTA literature, which documents formalization, poverty reduction, and firm entry following the trade agreement ([McCaig, 2011](#); [McCaig and Pavcnik, 2018](#); [McCaig et al., 2025](#)). Recent work shows that the same shock also affected family and human-capital margins. [McCaig et al. \(2022\)](#) show that the BTA affected youth migration, enrollment, work, and education spending, while [Mitra et al. \(2026\)](#) show that the BTA increased educational attainment and occupational mobility for both sons and daughters, with especially large effects on daughters' college education.² These findings provide empirical context for the quantity-quality channel in our model: export expansion changed children's human-capital margins, especially for daughters. We build on this evidence but study a different margin and use a different design. Rather than assigning the same provincial shock to every household, we use parent-specific industry-level exposure that separates maternal, paternal, and provincial channels within the same trade shock.

The results suggest that policies aimed only at raising household income are unlikely to break the link between better jobs for mothers and sex selection.³ The model points instead to policies that reduce the maternal time cost of children, such as institutional childcare and related supports. The education-price extension explains how per-child investments can amplify the mechanism, but it does not imply that lowering education prices alone would necessarily reduce sex selection: education costs enter both the price of fertility and the relative price of selecting sons. The cleaner policy implication is to loosen the maternal childcare constraint that makes higher female wages translate into lower fertility. The remainder of the paper presents the model, the empirical design, the main estimates, and tests that discriminate among alternative mechanisms.

2. Model

This section formalizes a household model with two margins, the maternal time cost of fertility and quantity-quality investment, and derives discriminating predictions for maternal, paternal, and provincial exposure. The baseline specification develops the fertility margin under maternal childcare time costs and son preference, and derives the shadow prices and comparative statics

²[Huynh and Ku \(2025\)](#) show, using province-level shift-share exposure, that the BTA reduced female labor force participation among married and wealthier women. [Drozdoff et al. \(2025\)](#) study how foreign manufacturing jobs in Vietnam delay marriage and childbirth through migration. We differ from both papers in shock, margin, or design, focusing on the BTA tariff shock and on sex composition at birth.

³[Anukriti \(2018\)](#) studies Devi Rupak, a Haryana cash-transfer program rewarding small families, especially those with daughters. Birth rates fell, but the sex ratio at birth worsened, consistent with the fertility-decline channel in our framework. Our paper differs in the source of the shock, a labor-market change that raises maternal opportunity cost, and in our design, which separates maternal from paternal exposure to isolate the source of income.

with respect to the maternal wage, paternal wage, and non-labor income. The [Becker and Lewis \(1973\)](#) quantity-quality margin enters through education investment per child, which adds education costs to the price of fertility while also making daughters more valuable when maternal wages rise. The resulting predictions allow the empirical analysis to distinguish the maternal time-cost mechanism from income and relative-returns explanations for rising sex ratios.

2.1 Baseline Model

Consider a household with pre-existing son preference, rooted in patrilineal kinship norms where sons provide old-age support and continue the family line ([Belanger, 2002](#)). The expected number of boys is $B = bn$, where n is total fertility and $b \in [0, 1]$ is the probability that a birth is male, which the household can influence through sex-selective practices.⁴ The household chooses consumption c , the mother's leisure ℓ , total fertility n , and the male birth probability b to maximize

$$\max_{c, \ell, n, b} U(c, \ell, n, B) \quad \text{where } B = bn, \quad (1)$$

with U increasing and strictly concave in each argument.⁵

Each child requires γ units of the mother's time, which competes with market work and leisure.⁶ The mother allocates her unit time endowment among leisure, market work, and childcare:

$$\ell + h_m + \gamma n = 1, \quad (2)$$

where h_m is her labor hours.⁷ The mother earns wage w_m , and the household receives non-labor income I , such as income from land, transfers, and remittances. Substituting the time constraint (2), the budget constraint is

$$c + w_m \ell + w_m \gamma n + p_b b n = W, \quad (3)$$

where $W = I + w_m + w_d$ is full household income and p_b is the cost per expected boy from sex selection.⁸ The parameter p_b encompasses the monetary costs of sex-determination technology (ultrasound fees, sex-selective abortion) and indirect costs such as health risks to the mother, psychological distress, and social stigma.

The first-order conditions with respect to (c, ℓ, n, b) are derived in [Appendix A.1](#), which also

⁴The variable b translates directly to the sex ratio at birth: $\text{Sex Ratio} = 100 \times b / (1 - b)$. For example, $b = 0.52$ corresponds to 108 males per 100 females.

⁵Assume U is twice continuously differentiable and the solution is interior.

⁶In our data, Vietnamese women spend over five hours per day on childcare compared to men's three hours ([ActionAid-Vietnam, 2016](#)). Formal childcare services remain limited outside major urban centers, and social norms assign childcare primarily to mothers. This asymmetry matches the broader time-use regularity emphasized by [Doepke et al. \(2023\)](#): women bear most childcare time costs, and fertility is lower where childcare burdens fall more disproportionately on women.

⁷We normalize paternal labor supply to one, so the father contributes earnings w_d but bears no childcare time cost. This abstraction isolates the maternal time-cost channel and is consistent with evidence that men's labor supply is largely unresponsive to fertility ([Angrist and Evans, 1998](#)).

⁸The income-side budget is $c + p_b b n \leq I + w_m(1 - \ell - \gamma n) + w_d$. Rearranging gives the full-income form.

provides the baseline comparative statics:

$$\frac{U_\ell}{U_c} = w_m, \quad \frac{U_n + U_B b}{U_c} = w_m \gamma + p_b b, \quad \frac{U_B n}{U_c} = p_b n. \quad (4)$$

Using the sex-selection condition to simplify the fertility condition gives effective shadow prices $(1, w_m, w_m \gamma, p_b)$ for the four goods (c, ℓ, n, B) . The maternal wage enters two of the four prices. The paternal wage enters none, affecting demands only through full income W .

The price ratio between fertility and boys,

$$\frac{U_n}{U_B} = \frac{w_m \gamma}{p_b}, \quad (5)$$

is the key object for comparative statics. When the maternal wage rises, $w_m \gamma$ increases and children become more expensive relative to expected boys. The fertility-price channel reduces desired fertility and raises the desired boy share $b = B/n$: with son preference, a household that has fewer births has a stronger incentive to concentrate those births on sons. The higher leisure price, $p_\ell = w_m$, reinforces the same substitution away from maternal time-intensive activities. The income effect is weakly positive if the income elasticity of boys is at least as large as that of fertility. Together these forces imply $\partial b / \partial w_m > 0$. For fertility, the substitution effect is unambiguously negative, and when it dominates the income effect, $\partial n / \partial w_m < 0$. Maternal labor supply rises because a higher wage draws the mother into market work and lower fertility frees up childcare time.

The paternal wage w_d operates through income alone, since it enters no effective price. Under the same income-elasticity condition, the income effect on b is weakly positive, but it lacks the substitution force that drives the maternal channel. The model therefore predicts at most a small positive paternal effect on sex selection, operating only through the income effect. The paternal wage raises fertility ($\partial n / \partial w_d > 0$) through income. Non-labor income I operates through the same income-only channel.

The baseline model connects panels (b), (c), and (a) of Figure 1: rising female wage employment raises the maternal time cost of children, lowers fertility, and increases the male share at birth. The remaining trend is the closing education gender gap in panel (d). The baseline treats each child's cost as maternal time only, but in practice wages also determine how much parents invest in children. The following extension adds this quantity-quality margin.

2.2 Human Capital Investment

The intuition for this extension is as follows. Because daughters' labor market returns also rise with the maternal wage, parents increase investment in daughters especially, narrowing the son-daughter education gap and raising the daughter-education component of child costs. The higher per-child cost reinforces the time-cost pressure to reduce fertility. At the same time, when daughters' investment catches up, replacing a daughter with a son adds less extra education spending

than before. Both price channels push toward a higher male birth share. Higher daughter returns also make daughters more valuable, partially offsetting the incentive to select for sons. The asymmetry with the paternal side is equally important. When paternal wages rise, sons' returns increase and parents invest more in sons. The resulting son-daughter investment gap raises the cost of selecting a son relative to a daughter, offsetting the income and son-value forces that would otherwise raise sex selection.

Parental wages serve a dual role. Beyond determining time costs, they signal the labor market returns to children's human capital. Children's labor market outcomes correlate with same-sex parental earnings, so maternal wages proxy for the return on daughters' education and paternal wages proxy for the return on sons' education. We formalize this by setting $r_g = w_m$ and $r_b = w_d$, where r_g and r_b denote the returns to girls' and boys' human capital.⁹ This extension embeds our model in the quantity-quality framework of [Becker and Lewis \(1973\)](#), where the total cost of children's education scales with both the number of children and the investment per child.

Parents invest e_b in each son and e_g in each daughter, where e measures educational input (such as years of schooling). A human capital production function $h(e)$ converts educational input into effective units of human capital, with $h' > 0$, $h'' < 0$, and $h(0) = 0$. The market return r prices each efficiency unit, so a child's human capital value is

$$H_b = r_b \cdot h(e_b), \quad H_g = r_g \cdot h(e_g). \quad (6)$$

Parents value children's human capital through $V(\bar{E})$ appended to the baseline utility, where $\bar{E} = b r_b h(e_b) + (1 - b) r_g h(e_g)$ is the expected per-child human capital and V is increasing and concave. The household now chooses $(c, \ell, n, b, e_b, e_g)$ jointly to maximize

$$U(c, \ell, n, B) + V(\bar{E}) \quad (7)$$

subject to the budget constraint

$$c + w_m \ell + n [w_m \gamma + p_e (b e_b + (1 - b) e_g)] + p_b b n = W, \quad (8)$$

where p_e is the unit cost of educational input (tuition, school fees). The key feature of this budget constraint is that education costs scale with both the number of children and investment per child. Adding one more child means paying for that child's education at the chosen quality level. Raising education investment means paying the additional schooling cost for every child. Each margin raises the cost of the other, so parents trade off between quantity and quality.

⁹The mapping $r_g = w_m$ is most natural for individual-level variation, where a mother's own industry determines both her wage and the labor market signal for daughters' future earnings. At the province level, returns may differ from individual wages when local industry composition shifts daughters' future prospects independently of the mother's own wage. Appendix [A.2.1](#) generalizes the model to allow returns to differ from individual wages, and Section [7](#) tests this empirically using female-weighted provincial exposure P_p^w .

Sex composition matters because sons and daughters may require different education spending. When $r_b > r_g$, sons receive more investment than daughters, so replacing a daughter with a son raises costs by $p_e(e_b - e_g)$. Fertility and quality are also linked: when fertility falls, parents invest more in each remaining child, raising the marginal cost of another child. This quantity-quality tradeoff connects education investment to both fertility and sex selection.

The first-order conditions for e_b and e_g (Appendix A.2) equate the marginal return per dollar across genders:

$$r_b h'(e_b^*) = r_g h'(e_g^*). \quad (9)$$

Because h is concave ($h'' < 0$), the gender with higher market returns receives more investment. When $r_b > r_g$, $e_b^* > e_g^*$. When maternal wages rise relative to paternal wages, daughters' returns increase and the investment gap narrows, matching the closing education gender gap in Figure 1, panel (d).

Grouping terms by n and $B = bn$ decomposes the budget into a fertility component and a sex-selection component:

$$c + w_m \ell + n[w_m \gamma + p_e e_g] + B[p_b + p_e(e_b - e_g)] = W. \quad (10)$$

The fertility price $p_n = w_m \gamma + p_e e_g$ now includes both childcare time and daughter education costs. The sex-selection price $p_B = p_b + p_e(e_b - e_g)$ reflects the incremental cost of a son relative to a daughter. When $w_d > w_m$, sons receive more investment than daughters ($e_b^* > e_g^*$), so the education gap $p_e(e_b - e_g) > 0$ adds to the cost of sex selection beyond p_b . Reducing γ directly weakens the maternal time-cost channel. Changing p_e is not an equivalent policy lever, because p_e enters both p_n and p_B and the two movements work against each other. Appendix A.2 provides further discussion on this asymmetry.

Consider the effect of a rise in w_m on sex selection. The human capital extension introduces two channels that work in opposite directions. A cost channel operates through the effective prices. Higher $r_g = w_m$ raises daughters' education investment, which increases the fertility price p_n and, when the education gap narrows, lowers the sex-selection price p_B . Both movements amplify the baseline time-cost effect. By contrast, a value channel operates through $V(\bar{E})$. Higher daughters' earnings make them more valuable to the household, reducing the incentive to select for sons.

The cost channel pushes b up through multiple reinforcing price effects: the higher leisure price, the amplified fertility price, and the lower education-cost penalty for selecting sons when daughters' investment catches up. The income effect reinforces the same direction, while only the value channel pushes b down. The net maternal effect on sex selection is likely positive.

This extension also changes the paternal wage prediction. Higher w_d raises sons' returns ($r_b = w_d$), making sons more expensive to educate. The higher cost of educating sons raises the sex-selection price p_B and offsets the income and son-value forces that push toward more sex selection. Because w_d does not raise the maternal time price of children, these opposing forces can

cancel, yielding an ambiguous and likely small paternal effect on sex selection. For fertility, the leading paternal force remains the positive income effect; any increase in p_n from higher desired child quality is indirect and second-order. Non-labor income I enters only through full income W without affecting education returns, leaving only the income effect.

Thus the extension adds panel (d) of Figure 1 to the baseline account: maternal wage gains can narrow the education gender gap while still reinforcing the pressure toward fewer births and a more male-skewed birth ratio.

These predictions apply to settings like Vietnam where fertility can adjust. Under binding fertility constraints, as under China’s one-child policy, the model generates different predictions (Appendix A.2.2 provides the details). When the fertility margin is constrained, the fertility-price channel is shut off, so the marginal maternal-wage effect works through leisure prices, sex-selection prices, income, and the value of daughters rather than through a reduction in n . Separately, income gains that have only income-effect responses in our setting may become positive predictors of sex selection, because the composition of births becomes the main remaining adjustment margin. At the same time, high per-child investment (a consequence of the binding quantity constraint) amplifies the value channel. When female returns rise, daughters become substantially more valuable because each daughter receives heavy investment. The relative-returns channel, which is dominated by time costs in Vietnam, gains force when the fertility margin is shut down. Thus the same framework highlights maternal time costs in Vietnam, while also explaining why income shocks and female returns can have stronger composition effects in China’s constrained-fertility setting (Almond et al., 2019; Qian, 2008).

Testable predictions. Table 1 collects the comparative statics. The rows distinguish maternal wages, paternal wages, and non-labor income, while the columns show the predicted signs for sex selection, maternal labor supply, and fertility. We highlight three predictions that organize the empirical analysis.

Prediction 1 (Maternal time cost). A rise in maternal wages increases male birth probability ($\partial b / \partial w_m > 0$), raises maternal labor supply, and reduces fertility. The maternal time constraint drives this pattern, while the quantity-quality margin reinforces the fertility and sex-selection responses.

Prediction 2 (Paternal income). A rise in paternal wages has at most weak effects on sex selection, because the education-cost channel offsets the income and son-value forces. The clearest observable response is higher fertility ($\partial n / \partial w_d > 0$) through the household income effect rather than the maternal time constraint.

Prediction 3 (Non-labor income). Non-labor income enters only through the income effect. It can raise fertility while leaving the male birth share nearly unchanged when expected boys and total fertility have similar income elasticities.

The key empirical test is whether this cross-channel pattern appears in the data: a maternal re-

Table 1: Model Predictions of How Income Sources Affect Household Outcomes

Increase in:	Outcome		
	(1)	(2)	(3)
	Sex Ratio b (male share)	Maternal Labor h_m	Fertility n
Mother's wage (w_m)	+	+	−
Father's wage (w_d)	0	−	+
Non-labor income (I)	0	−	+

Notes: + positive, − negative, 0 small or negligible predicted effect. Predictions summarize the model's key comparative statics in Section 2; Appendix A provides the derivations.

sponse across sex selection, labor supply, and fertility, without the same combination for paternal or provincial exposure. Section 7 uses this pattern to evaluate alternative explanations.

3. Data and Measures

We bring the theoretical framework to the data using Vietnam's 2001 Bilateral Trade Agreement (BTA) shock. This section constructs three exposure measures corresponding to the maternal-wage, paternal-wage, and broad local-income channels summarized in Table 1, then describes the household survey and census data used to observe parents' industries and birth outcomes.

3.1 The 2001 BTA and Exposure Measures

Before the BTA went into effect on December 10, 2001, Vietnamese exports to the US faced high tariffs under the Column 2 schedule applied to non-market economies. The BTA granted Vietnam Most Favored Nation (MFN) status, dramatically reducing US tariffs on Vietnamese exports while having minimal impact on Vietnam's import tariffs.¹⁰ Following the agreement, Vietnam's exports to the US surged from 0.7 billion dollars in 2000 to nearly 10 billion dollars in 2008.

We construct two complementary measures to capture how these differential tariff cuts transmit to fertility and sex selection through the labor market channels described in the model. The first captures direct exposure through parents' industries:

$$\tau_j = \log(1 + \text{Column } 2_j) - \log(1 + \text{MFN}_j), \quad (11)$$

where j denotes a parent's industry. A higher value of τ_j indicates a larger effective industry tariff cut. Because mothers and fathers tend to work in different industries, we construct separate

¹⁰While Vietnam was obligated to reduce tariffs mainly on agricultural and food products, these reductions were minimal, ranging from 0.03 to 2.7 percentage points. In addition to tariffs, Vietnam agreed to eliminate various import quotas. Most of these quotas had, however, been eliminated by the end of 2002 (STAR-Vietnam, 2003).

measures for each parent to capture relative changes in labor market returns.¹¹ These industry-level measures proxy for the maternal and paternal wage channels (w_m and w_d) in Table 1. Prior work using the same tariff variation shows that the BTA generated meaningful labor-market adjustment: [McCaig and Pavcnik \(2018\)](#) document large-scale reallocation from informal to formal employment, and [McCaig et al. \(2025\)](#) show growth in foreign and domestic firm entry in more exposed industries. We confirm this formalization channel directly in our own sample of parents at the time of fertility decisions below. Because female labor force participation was already approximately 70 percent before the BTA, the shock’s labor market effects on employed mothers operated primarily through wages and hours rather than through new labor force entry.

The second measure is a province-level shift-share that captures broader local economic effects:

$$\tau_p = \sum_j s_{jp} \tau_j, \quad (12)$$


where s_{jp} is the share of province p ’s employment in tradeable industry j in the 1999 Census. Following [Kovak \(2013\)](#), employment shares are computed over all tradeable industries. This provincial measure captures broad local economic effects rather than any individual household’s non-labor income. It proxies for the I channel in Table 1, but at the province level. Export expansion can raise returns to land and agricultural assets, and income gains can spread through transfers within extended families. Consistent with this broad local-income interpretation, [McCaig \(2011\)](#) use this province-level exposure measure to show that the BTA significantly reduced poverty rates across Vietnamese provinces. The provincial measure τ_p therefore captures reduced-form local gains that may flow through household and community networks rather than through a parent’s own industry alone.

Figure 2 illustrates the variation in both measures. Panel (a) shows the distribution of US tariff rates before and after the BTA across two-digit industries. The tariff reductions were large. The median Column 2 rate was 17.3 percentage points, falling to a median MFN rate of just 1.2 percentage points, a near-complete elimination of tariff protection. In log-difference terms, the industry tariff cut τ_j has a median of 0.168, reaching a maximum of 0.438 in the most exposed sector. A parent in the most exposed industry faced a tariff reduction roughly 2.6 times the median, implying substantially different incentives for formal employment and wage growth across the treatment distribution.

Panel (b) maps the resulting province-level exposure. The measure varies by a factor of 3.6 across 61 provinces, from 0.067 in the least exposed to 0.239 in the most exposed (mean 0.089, median 0.081). The geographic pattern is economically intuitive. Above-median exposure is concentrated in southern provinces around Ho Chi Minh City and along the central and southern

¹¹The correlation between maternal and paternal tariff exposure in our sample is consistently low, ranging from 0.24 to 0.36 across household survey years, described below, confirming that parents face different levels of exposure. The correlation between individual parental exposure and the provincial shift-share τ_p is even lower, remaining below 0.15 throughout the sample period. Appendix Figure C1 displays the full pairwise correlation structure by wave.

Figure 2: Variation in Trade Exposure Across Industries and Provinces



Notes: Panel (a) shows US tariff rates on Vietnamese exports before (Column 2) and after (MFN) the BTA across 2-digit ISIC industries, from [McCaig \(2011\)](#). Panel (b) shades provinces by shift-share exposure $\tau_p = \sum_j s_{jp} \tau_j$, where s_{jp} is industry j 's employment share in province p (1999 Census). Darker shading indicates above-median exposure.

coast, areas with established export infrastructure and a higher pre-existing share of employment in manufacturing. Northern inland and mountainous provinces, where employment was concentrated in agriculture, show the lowest exposure.

For the main analysis, we convert these continuous measures into binary treatment indicators by splitting at their sample medians. We define $M_i = \mathbf{1}[\tau_{j(m)} > \tilde{\tau}]$, equal to one if the mother's industry tariff cut exceeds the sample median $\tilde{\tau}$, and zero otherwise, and define F_i analogously for the father. We define $P_p = \mathbf{1}[\tau_p > \tilde{\tau}_p]$, equal to one for provinces with above-median shift-share exposure.¹²

We use binary indicators as the main specification because they deliver a transparent high-versus-low exposure contrast and avoid the weighting problems that arise in two-way fixed effects models with continuous treatments ([Callaway et al., 2024](#)). Section 6 reports continuous-treatment specifications as robustness checks.

3.2 Household Surveys and Census Data

To study how trade exposure affects sex ratio at birth, the ideal dataset would record each child's sex alongside both parents' industry, employment, and location around conception, allowing us to assign the trade shock separately to mothers, fathers, and provinces. The sex-selection literature

¹²Roughly 21 percent of parents work in industries where the Column 2 and MFN rates were both zero, yielding $\tau_j = 0$. These workers, concentrated in non-export service and agricultural sectors, fall in the low-treatment group.

typically relies on Demographic and Health Surveys (DHS) or population censuses, but DHS data rarely measure fathers' labor-market exposure, while infrequent censuses can observe parents' industry and location years after conception.¹³

We therefore use the Vietnam Household Living Standards Surveys (VHLSS) ([General Statistics Office of Vietnam, 2010](#)), a biennial nationally representative survey spanning 2002 through 2010 that records industry, location, and employment for both parents at each survey wave. Because the VHLSS does not contain complete birth histories, we work with the flow of births in each wave, treating each cross-section as a sample of infants aged 0–1 observed alive and co-resident with their parents. This choice trades off birth-history completeness for contemporaneous measurement of both parents' industries, which is essential for separating maternal and paternal exposure. We supplement the VHLSS with the 1999 Population Census from IPUMS International ([Ruggles et al., 2025](#)), whose province-level employment shares serve as weights for the shift-share exposure measure and whose birth records extend the pre-treatment baseline.

The Census provides births from 1997–1999, contributing approximately 49,000 infant observations to the pre-treatment baseline. The VHLSS waves from 2002 through 2010 contribute repeated cross-sections of infants aged 0–1, yielding approximately 17,700 infant-parent observations across five waves. For each infant we observe the child's sex and the parents' contemporaneous industry codes. In the VHLSS, we also observe parental work hours and job types. The Census therefore contributes to the sex selection analysis, while the labor supply and fertility analyses use VHLSS measures described below.

A key measurement challenge in this setting is identifying parent-child relationships within extended family households, where multiple generations often live together and survey relationships are recorded relative to the household head rather than within nuclear families. We develop an algorithm that identifies both couple and parent-child relationships within extended households, validated against known parent-child links in the 2014 and 2016 VHLSS waves, achieving 93 percent accuracy for mothers and 91 percent for fathers. [Appendix B](#) details this procedure.

[Table 2](#) presents summary statistics for Census 1999 and the five VHLSS waves. Mothers' baseline characteristics are comparable across sources, with average ages around 27–28 years and education ranging from 6.5 to 7.4 years across waves. The proportion of urban residents rises over time, consistent with Vietnam's urbanization trend. The male birth share is higher in the VHLSS waves than in the 1997–1999 Census baseline: the Census implies a sex ratio at birth of about 107 boys per 100 girls, while the 2002–2008 VHLSS waves imply ratios between 110 and 114. The estimation samples therefore capture the same broad upward shift documented in [Figure 1](#).

The model predicts that maternal and paternal wages have opposing effects on fertility ([Table 1](#), column 3). Our main analysis conditions on the birth of an infant aged 0–1, so fertility,

¹³DHS surveys provide birth histories, but they contain limited information on fathers and are conducted only intermittently. Population censuses are even more infrequent and come with a different problem: treatment must be assigned using industry and location at the census date rather than at conception. If workers switch industries or move after the shock, treatment is mismeasured and estimated effects are attenuated toward zero.

Table 2: Summary Statistics for Census 1999 and VHLSS Survey Waves

	Census	VHLSS				
	1999	2002	2004	2006	2008	2010
Male birth	0.518 (0.500)	0.532 (0.499)	0.529 (0.499)	0.524 (0.499)	0.527 (0.499)	0.511 (0.500)
Mother's age	27.00 (5.05)	28.10 (5.19)	28.39 (5.29)	27.88 (5.27)	28.09 (5.13)	28.24 (5.04)
Mother's educ. (yrs)	6.46 (3.43)	6.72 (3.60)	7.17 (3.71)	7.42 (3.69)	8.01 (3.71)	8.22 (3.66)
Ethnic minority	0.201 (0.401)	0.206 (0.404)	0.218 (0.413)	0.215 (0.411)	0.204 (0.403)	0.213 (0.409)
Urban	0.131 (0.337)	0.175 (0.380)	0.211 (0.408)	0.231 (0.421)	0.266 (0.442)	0.270 (0.444)
Birth order	1.93 (0.81)	2.07 (0.80)	2.02 (0.80)	1.90 (0.80)	1.88 (0.78)	1.89 (0.76)
Mother wage earner	—	0.202 (0.402)	0.243 (0.429)	0.275 (0.446)	0.333 (0.471)	0.366 (0.482)
Father wage earner	—	0.372 (0.483)	0.397 (0.489)	0.420 (0.494)	0.486 (0.500)	0.512 (0.500)
<i>N</i>	48,897	5,934	3,159	3,034	3,164	2,374

Notes: Standard deviations in parentheses. Census 1999 covers 1997–1999 births. VHLSS waves include children aged 0–1. Dashes indicate variables unavailable in the Census.

whether to have a birth at all, must be studied in a separate sample that includes woman-years without births. In the absence of complete birth histories, a common data challenge in this literature, we reconstruct annual birth outcomes from the household roster. Because the roster records the birth years of all children currently living with the mother, we can infer in which years she had a birth. For each woman in a given survey year, we create one observation per year from when she turned 20 until the survey date, coding the outcome as one if a child in the roster has the corresponding birth year. This reconstruction yields approximately two million woman-year observations for the fertility analysis.

The main limitation of the retrospective measure is that children who have left the household through marriage, migration, non-co-residence, or death are absent from the roster, so earlier births are systematically missing for older women. The measure also cannot observe fetal losses. We therefore interpret the fertility analysis as an annual birth-hazard test of whether the same

exposure variation changes the probability of a birth in a given year.

For the individual exposure measures, migration across provinces is not the direct source of treatment mismeasurement because maternal and paternal exposure are assigned from each parent’s industry, not location. The main concern is industry switching. Workers may move across industries in response to the trade shock, so that their current industry in the cross-section reflects post-shock sorting rather than pre-shock exposure. To address this selection problem, we exploit the VHLSS’s rotating panel feature, which tracks a subset of households across three consecutive waves (2002–2004–2006). From these households we draw women aged 20–40 as our analysis sample. By fixing each parent’s industry affiliation at the 2002 wave, the earliest available survey and effectively a pre-BTA baseline given the short window since December 2001 implementation, we isolate trade exposure from compositional changes driven by post-shock worker reallocation. The panel structure also allows us to control for time-invariant unobservables, including fertility preferences and son preference, through individual fixed effects. Given the short panel length, we rely on the cross-sectional variation for our main analysis and use the panel estimates as a robustness check.

Appendix Table C1 compares the 2002 wave of the rotating-panel sex-selection sample to the main 2002 cross-section on baseline demographics. Mothers’ age, education, minority status, urban share, and province-level tariff exposure are closely matched across the two samples, suggesting that the rotating-panel entrants are broadly similar to the main cross-section on observables.

4. Empirical Strategy

Testing the comparative statics in Table 1 requires an empirical design that separates direct maternal exposure, direct paternal exposure, and broader local exposure from a single trade shock. We embed the binary treatment indicators M_i , F_i , and P_p defined in Section 3.1 in a difference-in-differences framework that exploits industry-level tariff cuts and the province-level exposure induced by pre-BTA industry composition. Estimating the three channels jointly lets us distinguish the maternal labor-market mechanism from paternal-income and broader local-income channels within the same trade shock.

The pre- and post-treatment periods differ across outcomes. For sex selection and fertility, we set $\text{Post}_t = \mathbf{1}[\text{birthyear} > 2001]$, so that Census births from 1997–1999 are entirely pre-treatment while children born in 2002 or later are post-treatment. Since the BTA took effect in December 2001, the earliest post-treatment birth cohorts include children conceived before the agreement and are therefore only partially exposed. For maternal labor supply, we set $\text{Post}_s = \mathbf{1}[\text{survey year} > 2002]$, making the 2002 VHLSS the pre-treatment baseline and waves from 2004 onward post-treatment.

Our baseline specification for sex selection estimates a linear probability model where the unit

of observation is a child:

$$y_{it} = \beta_1(M_i \times \text{Post}_t) + \beta_2(F_i \times \text{Post}_t) + \beta_3(P_{p(i)} \times \text{Post}_t) + \alpha_t + \delta_{p(i)} + \gamma_{j_{m(i)}} + \gamma_{j_{d(i)}} + X'_{it}\phi + \varepsilon_{it}. \quad (13)$$

Here y_{it} is an indicator equal to one if child i born in year t is male, with $p(i)$, $m(i)$, and $d(i)$ denoting child i 's province of birth, mother, and father, respectively. The treatment indicators M_i , F_i , and $P_{p(i)}$ are assigned through the mother's industry, the father's industry, and the province of birth, respectively. The specification includes birth-year fixed effects α_t , province fixed effects $\delta_{p(i)}$, and ISIC industry fixed effects for the mother's industry $\gamma_{j_{m(i)}}$ and the father's industry $\gamma_{j_{d(i)}}$. The control vector X'_{it} includes the mother's education, urban residence, and minority status. It also includes world-average tariff controls at the parents' 2-digit ISIC industries matched by birth year, which absorb other industry-level tariff movements during the study period, including contemporaneous multilateral liberalization around Vietnam's WTO accession in January 2007. Interactions of pre-BTA provincial employment shares in agriculture and manufacturing with year fixed effects absorb differential trends arising from preexisting sectoral composition. Standard errors are clustered at the province level (61 clusters); Section 6 reports robustness to clustering at the industry level and two-way clustering.

For maternal labor market outcomes, the unit of observation shifts to a mother-survey wave. We estimate by OLS:

$$y_{ms} = \beta_1(M_m \times \text{Post}_s) + \beta_2(F_m \times \text{Post}_s) + \beta_3(P_{p(m)} \times \text{Post}_s) + \alpha_s + \delta_{p(m)} + \gamma_{j_m} + \gamma_{j_{d(m)}} + X'_{ms}\phi + \varepsilon_{ms}, \quad (14)$$

where y_{ms} is a labor market outcome for mother m observed in survey wave s , with $p(m)$ her province of residence and $d(m)$ her spouse, and $\text{Post}_s = \mathbf{1}[\text{survey year} > 2002]$. All remaining controls, fixed effects, and clustering follow Equation (13), with tariff controls matched by survey year rather than birth year. The coefficients β_1 , β_2 , and β_3 capture the post-BTA differential between high- and low-exposure groups.

For fertility, the specification replaces province and industry fixed effects with mother fixed effects. Using the retrospective birth panel described in Section 3.2, we define AnyBirth_{mt} as an indicator equal to one if mother m had a birth in year t :

$$\text{AnyBirth}_{mt} = \beta_1(M_m \times \text{Post}_t) + \beta_2(F_m \times \text{Post}_t) + \beta_3(P_{p(m)} \times \text{Post}_t) + \alpha_t + \mu_m + X'_{mt}\phi + \varepsilon_{mt}. \quad (15)$$

Mother fixed effects μ_m absorb all time-invariant heterogeneity, including underlying son preference, desired fertility, and permanent differences in industry and location. The control vector X'_{mt} includes interactions of pre-BTA provincial employment shares in agriculture and manufacturing with year fixed effects, and world-average tariff controls matched by birth year. Identification comes from within-mother variation in birth propensity before and after the BTA, comparing mothers with different levels of exposure.

Two threats to the cross-sectional estimates motivate the panel specification. The primary concern is industry switching. The cross-section records each parent’s current industry rather than pre-shock affiliation, so workers who move into high-tariff-cut industries after the BTA can contaminate M_i and F_i with post-shock selection. This concern is most relevant for the maternal and paternal industry coefficients, because treatment is assigned from each parent’s own industry. Migration across provinces is a secondary concern for the province-level coefficient. If the trade shock draws households with stronger son preference into high-exposure provinces, the composition of the high- P_p group changes in a way that biases β_3 upward. Including P_p alongside the parent-specific measures helps separate this broader province-level channel from maternal and paternal industry exposure, but it does not address sorting across industries.

To address these concerns, we exploit the VHLSS rotating panel described in Section 3.2 and lock each parent’s treatment status to her pre-BTA industry affiliation, following [McCaig and Pavcnik \(2018\)](#):

$$y_{it} = \beta_1(M_i \times \text{Post}_t) + \beta_2(F_i \times \text{Post}_t) + \beta_3(P_{p(i)} \times \text{Post}_t) + \alpha_t + \mu_{m(i)} + X'_{it}\phi + \varepsilon_{it}. \quad (16)$$

The unit of observation remains a child i , as in Equation (13), but mother fixed effects $\mu_{m(i)}$ replace province and industry fixed effects, absorbing all time-invariant heterogeneity of the mother $m(i)$ of child i . Treatment indicators M_i and F_i are assigned based on each parent’s industry reported in the 2002 VHLSS wave, the earliest available survey. Since the 2002 wave was fielded shortly after the BTA took effect, fixing exposure to this initial wave limits the scope for later industry switching to determine treatment status.

The panel sample consists of women aged 20–40 observed in the VHLSS rotating panel. For sex selection, the analysis uses mother-wave observations with eligible recent births and fixes treatment to each parent’s 2002 industry. Identification with mother fixed effects comes from comparing births to the same mother before and after the BTA when such within-mother variation is available. If endogenous sorting were inflating the cross-sectional estimates, fixing treatment to pre-shock industries should attenuate the panel coefficients relative to the cross-sectional results.

The key identifying assumption is that, absent the BTA, outcomes would have evolved similarly across high- and low-exposure groups. The 1999 Census births (1997–1999) and VHLSS births in 2000–2001 provide pre-treatment observations spanning multiple years before the shock, which we use to assess differential pre-trends in Section 5. For the panel specification, fixing treatment to the first VHLSS wave reduces the concern that later industry switching drives exposure, but it still requires that initial industry affiliation was not chosen in anticipation of the BTA’s differential tariff effects across sectors.

The coefficients β_1 , β_2 , and β_3 map directly to the net comparative statics in Table 1. The model predicts $\beta_1 > 0$ for sex selection and maternal labor supply and $\beta_1 < 0$ for fertility: maternal exposure raises the opportunity cost of childbearing, while the quantity-quality margin reinforces

the fertility and sex-selection responses. Paternal exposure (β_2) and province-level exposure (β_3) provide comparison channels because they shift household resources without the same maternal-time component.

5. Main Results

5.1 Effects on Sex Selection

Prediction 1 implies that sex selection should respond specifically to the mother's own labor-market exposure. Maternal exposure should raise male birth probability because it changes the maternal time cost of children and the quantity-quality margin, while paternal and province-level exposure should have little or no effect on sex selection. Table 3 reports progressively richer specifications of Equation (13), where the dependent variable is an indicator equal to one if the child is male.

Columns (1) and (2) introduce the maternal and paternal exposure channels individually, while column (3) combines them. Columns (4) and (5) add the province-level channel, with column (5) adding the full control set. When entered alone in column (1), maternal industry exposure ($M_i \times \text{Post}_i$) raises male birth probability by 6.5 percentage points (s.e. 0.025), significant at the 5 percent level. Against a dependent variable mean of 0.52, this corresponds to a 13 percent relative increase in the probability that a birth is male. The binary specification makes the channel comparison transparent, while the continuous treatment specification below scales the estimate to the size of observed tariff cuts.

For a mother whose industry received the sample median tariff cut of approximately 7 percentage points, the continuous treatment specification in Appendix Table C7 implies a 1.5 to 2.7 percentage point increase in male birth probability, depending on the control set. The raw means tell a similar story. The high-exposure group's male birth share rose from 51.4 to 53.3 percent after the BTA, while the low-exposure group rose from 52.1 to 52.5 percent, a difference-in-differences of 1.5 percentage points. At the lower end of the continuous-treatment range, this corresponds to an increase in the sex ratio at birth from about 108 to 114 boys per 100 girls, close to the aggregate movement shown in Figure 1.

The father's industry exposure ($F_i \times \text{Post}_i$) enters with a positive but imprecisely estimated coefficient of 0.04 in column (2). In the combined specifications of columns (3) through (5), the paternal coefficient falls to 0.02 and remains statistically indistinguishable from zero. This is the sex-selection implication of Prediction 2: paternal exposure raises household income without tightening the maternal time constraint, so it should have at most weak effects on male birth probability. Across these same specifications, the maternal coefficient moves from 0.065 in column (1) to 0.055 when the father is added in column (3) and increases modestly to 0.077 with the full control set in column (5). The full control set includes the world-average tariff controls described in Section 4.

Province-level exposure ($P_p \times \text{Post}_i$) is near zero and insignificant throughout, entering at 0.002 in column (4) and 0.011 in column (5). This result matches the comparison channel in Prediction 3.

Table 3: Effects of BTA Trade Exposure on Male Birth Probability

	MaleBirth				
	(1)	(2)	(3)	(4)	(5)
$M_i \times \text{Post}$	0.065** (0.025)		0.055** (0.025)	0.055** (0.024)	0.077** (0.033)
$F_i \times \text{Post}$		0.045 (0.029)	0.027 (0.029)	0.026 (0.029)	0.035 (0.029)
$P_p \times \text{Post}$				0.002 (0.012)	0.011 (0.015)
R^2	0.015	0.014	0.015	0.015	0.016
Observations	64,170	64,170	64,170	64,170	64,170
Dep. Var. Mean	0.518	0.518	0.518	0.518	0.518
Controls					✓


Notes: Dep. var. is MaleBirth. M_i , F_i , P_p = 1 if maternal, paternal, or provincial tariff cut exceeds the sample median. All columns include birth-year, province, and parents' 2-digit ISIC fixed effects. Controls (col. 5) add mother's education, urban residence, minority status, world-average tariff controls at parents' industries, and pre-BTA provincial employment shares $\times$ birth-year. Province-clustered standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Broad local income gains do not directly raise the mother's opportunity cost of childbearing, so they should have weaker effects on sex selection than direct maternal industry exposure.

Column (5) presents the fully controlled specification with all three exposure channels entering simultaneously. Across the progressive specifications, the maternal coefficient moves from 0.065 in the maternal-only column (1) to 0.077 in the fully controlled column (5), remaining significant at the 5 percent level throughout. To place the magnitude in context, [Anukriti \(2018\)](#) finds that first births become 0.5 to 1.2 percentage points more likely to be male after the Devi Rupak fertility incentive in India, and [Almond et al. \(2019\)](#) estimate a cumulative increase of approximately 2.8 percentage points over four years for second births following a firstborn girl in China. Our continuous-treatment estimate of 1.5 to 2.7 percentage points (Appendix Table C7) exceeds the upper bound of the range documented by [Anukriti \(2018\)](#) and falls below the benchmark from [Almond et al. \(2019\)](#), consistent with the fact that the latter conditions on the subgroup facing the strongest son preference pressure while our estimate pools across all birth orders and sibling compositions. The heterogeneity analysis below explores whether the maternal exposure effect varies across subgroups.

Figure 3 traces these effects over time, plotting event study estimates from the full specification where all three exposure channels enter simultaneously with the complete set of controls. Appendix Figure C2 presents individual event studies for each channel estimated separately, con-

Figure 3: Event Study Estimates of BTA Exposure on Male Birth Probability



Notes: Coefficients from a single regression of Equation (13) with all three exposure channels and full controls. Panel (a): maternal (M_i), panel (b): paternal (F_i), panel (c): provincial (P_p). Binary indicators split at the sample median. Omitted year is 2001. Error bars show 95% confidence intervals with province-clustered standard errors.

firming that the patterns are not artifacts of joint estimation.


Panel (a) displays the maternal exposure coefficients. Pre-BTA estimates for the 1997–1999 Census cohorts and the 2000–2001 VHLSS births cluster tightly around zero, confirming that high- and low-exposure groups followed parallel trends well before the shock.¹⁴ The 2002 coefficient is negative and imprecise. The maternal estimates turn positive in 2003 and 2004, with similar magnitudes and marginal significance, and become larger and statistically significant from 2005 onward. The timing is consistent with gestation: because the BTA took effect in December 2001, most births recorded in 2002 reflect conceptions that predate the agreement, and affected births would appear in later cohorts.

Panels (b) and (c) display paternal and province-level exposure. Neither channel exhibits a consistent post-BTA pattern. Paternal industry exposure fluctuates around zero in most years, with the exception of an isolated 2004 spike.¹⁵ Province-level exposure also lacks a persistent positive pattern, with estimates small and statistically insignificant at the 5 percent level. Pre-trend coefficients are flat for both channels, ruling out the concern that differential pre-existing trends could explain the maternal results through correlated shocks. The contrast across the three panels

¹⁴A reverse causality concern could arise because the design is cross-sectional. The mother’s industry is measured contemporaneously rather than before the birth, so a mother who had a son could in principle sort into a high-treatment industry afterward. Two pieces of evidence rule this out. First, the flat pre-BTA coefficients show no pre-existing differential trend between high- and low-exposure groups, ruling out pre-existing compositional differences correlated with treatment. Second, if having a son led parents to work more and sort into higher-wage industries, the same mechanism should produce a positive paternal coefficient, but panel (b) shows no such effect.

¹⁵Annual event-study cells are small, so we do not interpret isolated single-year movements on their own. The 2004 paternal estimate is not followed by a persistent positive pattern in adjacent years, unlike the maternal series.

Figure 4: Heterogeneous Effects of Maternal Industry Exposure on Male Birth Probability



Notes: Each point is the coefficient on $M_i \times \text{Post}_t$ from Equation (13) estimated within the indicated subgroup, with 95% confidence intervals. Pre-BTA SRB split at median male birth share among 1997–1998 Census births. All regressions include full controls and fixed effects.

is the central identification result. The maternal time-cost mechanism predicts $\beta_1 > 0$ with β_2 and β_3 small for sex selection (Table 1, column 1). The event study confirms that the positive pattern is unique to maternal exposure and emerges only after the BTA could plausibly have affected birth timing.

Because maternal exposure also reduces births, the fertility results below return to whether selection into realized births can explain the male-birth response.

Heterogeneity. We examine how the maternal exposure effect varies across subgroups by estimating Equation (13) separately within each group, with the same fixed effects and controls as the main specification. Figure 4 displays the coefficient on $M_i \times \text{Post}_t$ and its 95 percent confidence interval for each subgroup.

The maternal exposure effect is positive in both rural and urban areas, with point estimates of 0.072 and 0.093, respectively. The rural estimate is statistically significant ($p < 0.05$) while the urban estimate is imprecise, reflecting the smaller urban subsample. This pattern does not fit a purely urban technology-access explanation, which would predict a clearer concentration of effects in cities where sex-selective technology is more available. The significant rural response is consistent with an industry-level labor-demand shock operating where son preference norms may bind more strongly.

We split provinces by their pre-BTA sex ratio, using the median male birth share among 1997–

1998 Census births as the cutoff (Figure C3 maps the provincial variation). Provinces with higher pre-existing sex ratios show a larger maternal exposure effect (0.105, $p < 0.05$) than provinces with lower pre-BTA sex ratios (0.047, not significant). The pre-BTA sex ratio proxies for the intensity of pre-existing son preference across provinces. Where son preference is stronger, the marginal value of sex selection is higher, so the model predicts a larger response to a given increase in the maternal opportunity cost of childcare.

The parity split addresses the concern that the male-birth response reflects fertility selection based on prior child sex rather than prenatal sex selection. The first-birth estimate is positive and marginally significant (0.063, $p < 0.10$), which is informative because prior child sex cannot shape fertility choices before a first birth. The point estimate is larger at second parity (0.101) and remains positive at third and higher parities (0.081), though these estimates are imprecise.¹⁶ The higher-parity estimates point in the expected direction because parents at later parities have fewer remaining births over which to fulfill son preference.

Mothers with higher education (ten or more years of schooling) show a larger point estimate (0.139) than those with less schooling (0.047), with the high-education estimate marginally significant ($p < 0.10$). More educated mothers typically command higher wages in the labor market, so a given improvement in industry conditions implies a larger increase in their opportunity cost of time and a stronger incentive to economize on time-intensive childcare through sex selection. The estimates are too imprecise to make the education gradient conclusive, but the larger point estimate among more educated mothers is consistent with the time-cost interpretation.

Maternal age at birth generates the sharpest heterogeneity. Mothers aged 30 and above show a large and significant effect (0.144, $p < 0.05$), while mothers under 30 show no discernible response (0.033, not significant). Older mothers are more likely to be at higher parities without yet having a son, have accumulated work experience and face a higher opportunity cost of childcare, and confront a shorter remaining fertility horizon. Each of these factors reinforces the incentive to sex select and the pressure to produce a son before the fertility window closes.

Overall, the heterogeneity results point to a stronger response when maternal wage opportunities interact with stronger incentives to secure a son. The first-birth estimate is positive, and the largest responses appear in provinces with high pre-BTA sex ratios and among older mothers facing a narrowing fertility horizon. The next subsection tests the labor-market margin directly.

5.2 Effects on Parental Labor Market Outcomes

The sex-selection result requires a labor-market first stage: industry tariff cuts improve mothers' work opportunities. Prediction 1 then implies that maternal exposure should raise maternal labor supply. Table 4 tests these labor-market links by estimating the same specification on parental labor market outcomes. Columns (1)–(2) report effects on maternal wage employment, columns (3)–

¹⁶Births of order 3 or higher are uncommon in our data. Vietnam's total fertility rate had already fallen to 2.01 by 2000.

Table 4: Effects of BTA Exposures on Parental Labor Market Outcomes

	MomWageEmp		log(MomWorkHrs)		DadWageEmp	
	(1)	(2)	(3)	(4)	(5)	(6)
$M_i \times \text{Post}$	0.149*** (0.027)	0.149*** (0.027)	0.159*** (0.051)	0.155*** (0.051)	-0.015 (0.029)	-0.014 (0.029)
$F_i \times \text{Post}$	0.018 (0.019)	0.018 (0.018)	0.000 (0.036)	-0.005 (0.037)	0.147*** (0.025)	0.148*** (0.024)
$P_p \times \text{Post}$		-0.003 (0.024)		0.054 (0.045)		-0.010 (0.038)
R^2	0.604	0.604	0.254	0.254	0.530	0.530
Observations	17,389	17,389	17,198	17,198	17,389	17,389
Dep. Var. Mean	0.247	0.247	4.66	4.66	0.406	0.406
Controls	✓	✓	✓	✓	✓	✓

Notes: Dependent variables are maternal wage employment (cols. 1–2), log monthly hours (cols. 3–4), and paternal wage employment (cols. 5–6). Treatment indicators and fixed effects as in Table 3. Controls include education, urban, minority status, and world-average tariff controls at parents' industries (survey-year matched). Province-clustered standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

(4) on maternal log monthly hours, and columns (5)–(6) on paternal wage employment.

Column (1) shows that mothers in high-tariff-cut industries are 15 percentage points more likely to hold a wage job after the BTA ($p < 0.01$), from a baseline of 25 percent. The estimate is stable in column (2) when provincial exposure is added, and paternal and province-level exposure effects on maternal wage employment are small and statistically insignificant.

Wage employment is the relevant first-stage measure in this setting because, although female labor force participation was approximately 70 percent before the BTA, only one in four mothers held a wage job, leaving most working mothers in informal or agricultural work without an observed market wage. Estimating wage effects only among wage workers would condition on post-treatment selection into wage jobs, mixing true wage growth with changes in the composition of observed wage workers. The transition into wage employment from informal or agricultural work is itself a signal of improved labor market conditions, and the shift from the informal to the formal sector is precisely the reallocation that [McCaig and Pavcnik \(2018\)](#) document using the same tariff variation. In the Vietnamese context, formalization raises the value of market work through more stable wage jobs while also making work schedules less flexible, both of which increase the opportunity cost of time spent away from employment.

On the intensive margin, columns (3)–(4) show that mothers in high-exposure industries work 0.15 log-points more monthly hours ($p < 0.01$), approximately a 15 percent increase from a base-

Table 5: Effects of BTA Exposures on Fertility (Probability of Giving Birth)

	AnyBirth				
	(1)	(2)	(3)	(4)	(5)
$M_i \times \text{Post}$	-0.004*** (0.001)		-0.006*** (0.001)	-0.006*** (0.001)	-0.007*** (0.001)
$F_i \times \text{Post}$		0.003* (0.001)	0.005*** (0.001)	0.005*** (0.001)	0.005*** (0.001)
$P_p \times \text{Post}$				-0.0001 (0.002)	-0.002 (0.002)
R^2	0.29	0.29	0.29	0.29	0.29
Observations	2,011,963	2,011,963	2,011,963	2,011,963	2,011,963
Dep. Var. Mean	0.02	0.02	0.02	0.02	0.02
Controls					✓

Notes: Dependent variable is an indicator equal to one if the mother had any birth in the past 12 months. Sample: mothers aged 20–40 from the VHLSS 2002–2008 retrospective birth panel (approximately 2 million woman-year observations). All specifications include mother fixed effects and birth-year fixed effects. Controls in column (5) add world-average tariff controls at parents’ 2-digit ISIC industries matched by birth year, and interactions of pre-BTA provincial employment shares in agriculture and manufacturing with birth-year dummies. Province-clustered standard errors in parentheses (61 clusters). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

line of about 110 hours per month. Because the BTA shifts some mothers from informal or agricultural work into wage employment, the hours estimates combine changes in hours within existing jobs with movement into wage jobs that carry more regular time commitments. We therefore interpret the hours result as evidence of increased maternal market-work attachment rather than as a pure intensive-margin response for a fixed job.

Columns (5)–(6) show the mirror pattern for fathers. Maternal industry exposure ($M_i \times \text{Post}_t$) is near zero for paternal wage employment, while paternal industry exposure ($F_i \times \text{Post}_t$) enters at 15 percentage points ($p < 0.01$). The labor-market response is therefore parent-specific: each parent’s wage employment responds to exposure in that parent’s own industry. This result verifies the labor-market first stage behind Prediction 2, but only the maternal channel transmits to sex selection.

5.3 Effects on Fertility

Predictions 1 and 2 generate opposing fertility effects. Maternal exposure reduces births through time costs, while paternal exposure raises births through income (Table 1, column 3). Table 5 reports estimates where the dependent variable is an indicator equal to one if the mother had any birth in the past 12 months.

The retrospective woman-year panel follows Equation (15) with mother fixed effects. The estimates compare changes in a mother's annual birth propensity around the BTA across exposure groups. Maternal industry exposure reduces the probability of giving birth. Column (1) yields a coefficient of -0.004 ($p < 0.01$) when maternal exposure enters alone, and the coefficient remains negative and significant across the combined specifications in columns (3) through (5), reaching -0.007 in the full specification. The full-specification estimate corresponds to a 0.7 percentage-point decline in the annual probability of giving birth. Because this outcome is reconstructed from children observed in the household roster, its mean should be read as the mean of an annual birth-hazard measure, not as an estimate of the total fertility rate. Cumulated mechanically over a 20-year reproductive window, the estimated annual hazard effect would imply about 0.14 fewer births per woman, roughly 7 percent of a total fertility rate of 2. The sign is consistent with the model prediction that maternal industry exposure reduces fertility by raising the opportunity cost of the mother's time.

Paternal industry exposure moves in the opposite direction. Once maternal exposure is included, the paternal coefficient is 0.005 ($p < 0.01$) across columns (3) through (5). This estimate corresponds to a 0.5 percentage-point increase in the annual probability of giving birth in the reconstructed fertility panel. Cumulated mechanically over a 20-year reproductive window, this annual hazard effect would imply about 0.10 additional births per woman, roughly 5 percent of a total fertility rate of 2. The positive sign is consistent with Prediction 2: paternal exposure raises fertility through an income effect without tightening the maternal time constraint.

Province-level exposure ($P_p \times \text{Post}_t$) is near zero and statistically insignificant in columns (4) and (5). As with sex selection, broad local exposure does not reproduce the direct maternal pattern. Prediction 3 would imply a positive fertility response if this measure operated mainly as a pure income shock. The null result therefore suggests that our province-level exposure measure is not functioning as a strong pure-income channel once parent-specific exposure is included.

The opposing signs of maternal and paternal exposure on fertility are difficult to reconcile with a model where any positive income shock affects childbearing in the same direction. They instead point to the distinction between maternal time costs and household income. Maternal exposure raises the opportunity cost of childbearing and reduces births, while paternal exposure relaxes the household budget and raises births. The opposing signs also address the possibility that higher income could ease childbearing by allowing households to purchase childcare. If purchased childcare were a perfect substitute for the mother's own time, a maternal wage gain would look more like a pure income shock and should raise fertility, as paternal exposure does. This pattern is consistent with childcare remaining a binding maternal time cost. Vietnam's institutional context, in which formal childcare is scarce outside major urban centers and most working mothers rely on family or informal arrangements, reinforces this conclusion.

One qualification is that the 12-month birth indicator captures an annual birth hazard rather than completed fertility. Part of the estimated reduction may therefore reflect delayed births rather

than a permanent decrease in desired family size. Sex-selective abortion could also lengthen birth spacing, which would make the annual hazard effect larger than the effect on completed fertility. Appendix Table C11 varies the retrospective lookback window. The maternal coefficient is -0.003 when the reconstruction is limited to births within 10 years of the survey, then grows to -0.005 with a 12-year window and stabilizes around -0.006 to -0.007 for wider windows. The negative coefficient persists as we allow births farther back from the survey date, suggesting that the result is not an artifact of a particular retrospective window. We therefore interpret the fertility estimates as evidence that maternal exposure changed the timing and incidence of births in the direction predicted by the time-cost mechanism.

The sex-ratio estimates use births as the unit of observation, so they measure the male share among births that occur and are observed in the household. Because maternal exposure also reduces the annual birth hazard, a possible selection concern is that treatment changes which births remain in the sample. A selection-only explanation would require the births shifted out of the observed window to be disproportionately female, raising the male share among realized births even without a change in intended birth composition. Several patterns make this explanation unlikely to account for the full result. The maternal coefficient remains positive among first births, where prior child sex cannot shape subsequent fertility choices. Paternal exposure raises fertility but does not move the sex ratio. The event study shows no differential pre-trend, and the maternal effect appears only after the BTA could affect labor demand and conception timing. These patterns point toward sex selection among realized births, without fully decomposing prenatal selection from birth timing.

5.4 Panel Evidence

The cross-sectional results align with the model, but they rely on parent industries measured at the time of the survey. The panel evidence addresses the specific sorting concern that women with stronger son preference may have moved into high-tariff-cut industries after the BTA, inflating the maternal exposure coefficient. Table 6 re-estimates the sex-ratio specification using the VHLSS rotating panel and separates the design change into steps. Column (1) reproduces the maternal exposure estimate in the 2002–2006 cross-section. Column (2) restricts the sample to mothers observed in the rotating panel and fixes treatment using the mother’s 2002-wave industry. Column (3) replaces province and industry fixed effects with mother fixed effects, and column (4) adds the paternal and provincial exposure channels. The table therefore asks whether the maternal effect survives when exposure is fixed before later industry switching and identification comes from within-mother changes around the BTA.

The maternal coefficient is 0.06 ($p < 0.05$) in column (1), matching the full-sample estimate in Table 3. It rises to 0.17 ($p < 0.05$) in column (2), which uses the rotating-panel sample and assigns exposure from the mother’s 2002-wave industry. Among women working in 2002, approximately 23% changed their 2-digit ISIC industry by 2004 or 2006. Fixing treatment at the 2002 industry

Table 6: Panel Evidence on Within-Mother Identification

	MaleBirth			
	(1)	(2)	(3)	(4)
$M_i \times \text{Post}$	0.06** (0.03)	0.17** (0.07)	0.34*** (0.08)	0.30*** (0.11)
$F_i \times \text{Post}$				0.05 (0.16)
$P_p \times \text{Post}$				0.03 (0.06)
R^2	0.01	0.09	0.83	0.82
Observations	12,127	2,684	2,788	2,615
Mothers		1,723	1,723	1,723
Dep. Var. Mean	0.53	0.52	0.52	0.53
Individual FE			✓	✓
Industry FE	✓	✓		
Province FE	✓	✓		
Demographic controls	✓	✓	✓	✓

Notes: Column (1) uses the full 2002–2006 cross-section with contemporaneous industry assignment. Columns (2)–(4) restrict to 1,723 panel mothers with 2002-wave industry assignment. Columns (3)–(4) use mother fixed effects, with the within-mother comparison driven by mothers observed with births on both sides of the BTA. Column (4) adds paternal and provincial exposure channels. Demographic controls are mother’s education, urban residence, and minority status. Province-clustered standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

therefore reduces the concern that contemporaneous post-BTA occupation changes mechanically generate the effect. Table C1 shows that the 2002 rotating-panel baseline is broadly similar to the broader cross-section on education, ethnicity, urban residence, and province-level tariff exposure, with the main difference being that mothers in the panel are about one year older on average.

Columns (3) and (4) replace province and industry fixed effects with individual mother fixed effects, absorbing time-invariant maternal heterogeneity including baseline son preference and permanent industry sorting. In column (3), the maternal coefficient rises to 0.34 ($p < 0.01$). Although the regression includes 1,723 panel mothers, identification comes from the 312 mothers observed with births both before and after the BTA. Column (4) adds paternal and provincial exposure channels, and the maternal coefficient remains similar at 0.30 ($p < 0.01$). The paternal and provincial coefficients are individually small and statistically insignificant, consistent with the main results in Table 3. Because the within-mother sample is selected by repeated observed births, we use the panel estimate mainly as a check on post-BTA sorting rather than as the pre-

ferred specification.

Taken together, the cross-sectional and panel results trace a discriminating pattern across the three exposure channels and three outcomes that maps directly onto the model predictions in Table 1. The evidence matches the full maternal row: maternal industry exposure increases male birth probability, increases maternal labor supply, and reduces fertility, as Prediction 1 requires. The paternal results also fit the key sex-selection and fertility predictions: paternal exposure does not move the male birth share but raises fertility, consistent with Prediction 2. Provincial exposure matches the predicted null for sex selection but does not generate the positive fertility response predicted by Prediction 3 when local exposure operates as a pure non-labor income shock. The results therefore support the model most strongly where it is most distinctive, by showing that the sex-selection response appears only when the shock raises the mother’s own labor-market opportunity.

6. Robustness Checks

This section examines the sensitivity of our main results to alternative estimation choices. We consider alternative clustering schemes, region-by-year fixed effects, nonlinear specifications, and continuous treatment measures.

Alternative Clustering. Our baseline specifications cluster standard errors at the province level (61 clusters). Although tariff cuts vary by industry, the household-level treatment combines maternal, paternal, and provincial exposure. We therefore report industry and two-way clustering as robustness checks. Table C2 clusters by 59 2-digit ISIC industries, and Table C3 applies two-way clustering by province and industry to allow both geographic and sectoral correlation in the error term.

Under ISIC-only clustering (Table C2), all four maternal exposure effects remain statistically significant, with standard errors close to the baseline. Under two-way clustering (Table C3), the maternal effects also remain statistically significant, with standard errors close to the baseline for sex selection and labor supply and modestly larger for fertility. Across both approaches, the paternal exposure coefficients remain small and insignificant for sex selection and maternal labor outcomes, while the paternal fertility effect retains significance at the 5 percent level. The only coefficient outside the main maternal pattern that changes inference is provincial exposure for maternal hours, which is positive and marginally significant under alternative clustering. The model isolates non-labor income as a distinct theoretical channel, whereas the empirical province-level exposure measure is a broader reduced-form local shock. In practice, τ_p may also capture local labor-demand spillovers, including changes in sectors where women work, which could raise mothers’ hours. We therefore interpret this coefficient cautiously. It does not change the main conclusion that the sex-selection and fertility responses are tied to direct maternal industry exposure rather than broad provincial exposure.

Region-Year Fixed Effects. A potential concern is that regional infrastructure improvements could independently affect access to sex-selective technologies (Chen et al., 2013) or childcare services, confounding our estimates of trade exposure effects. Vietnam’s infrastructure development typically follows its eight administrative regions (Northeast, Northwest, Red River Delta, North Central Coast, South Central Coast, Central Highlands, Southeast, and Mekong River Delta), and shocks to healthcare access or childcare provision may vary across these regions over time. Since these regions correspond to official planning units, region-by-year fixed effects absorb region-level policy variation and differential investment in healthcare or childcare infrastructure. We incorporate region-by-year fixed effects into our baseline specifications to absorb any such time-varying regional confounders.

Table C5 shows that including region-by-year fixed effects leaves all four outcomes essentially unchanged. The maternal exposure coefficients retain their baseline magnitudes and significance levels across sex selection (0.08, $p < 0.05$), hours (0.15, $p < 0.01$), wage employment (0.15, $p < 0.01$), and fertility (-0.007 , $p < 0.01$). Paternal and provincial exposure measures remain small for sex selection and labor supply. The paternal fertility coefficient (0.005, $p < 0.01$) retains its significance, consistent with the main results. The results are unlikely to be driven by omitted regional shocks to infrastructure or healthcare access.

Nonlinear Estimation. Given the binary male-birth outcome, we check robustness to nonlinear specification using a fixed-effects logit model. Table C6 reports the logit estimates. The maternal exposure coefficient ($M_i \times \text{Post}$) is 0.22 without controls and 0.31 with controls, both significant at the 5 percent level. The implied odds ratio from the controlled specification is around 1.36, indicating a 36 percent increase in the odds of a male birth for mothers in high-exposure industries. The implied average marginal effect is close to the linear probability model estimate.

The paternal and provincial exposure coefficients remain small and statistically insignificant across both specifications, confirming that maternal industry exposure drives the sex-selection response. The agreement between logit and linear probability model results suggests that the main estimates are not an artifact of the linear functional form.

Continuous Treatment Exposures. Our main analysis uses binary indicators that classify workers as high-exposure or low-exposure based on whether their industry’s tariff cut exceeds the sample median. This binary split discards variation in exposure intensity within each group. We now replicate the analysis using continuous tariff-cut measures ($\tau_m = \tau_{j(m)}$, $\tau_d = \tau_{j(d)}$, and τ_p from Section 3.1), which exploit the full distribution of tariff reductions across industries. Tables C7, C8, and C9 present estimates for sex selection, maternal labor supply, and fertility.

The sex-selection result is robust to the continuous specification. Table C7 shows that the maternal exposure coefficient on male birth ($\tau_m \times \text{Post}$) is positive and significant in all specifications where it is included (columns 1 and 3–5). The coefficient ranges from 0.21 without controls to 0.37

with the full control set, with significance at least at the 5 percent level, paralleling the pattern in the binary specifications where adding controls strengthens the estimate. Column (2) isolates the paternal channel and does not include τ_m . Paternal and provincial exposure coefficients are small and insignificant, as in the binary results.

The maternal labor supply results in Table C8 confirm the wage-employment effect and remain directionally consistent on hours. The maternal exposure coefficient on wage employment is 0.51 without controls and 0.63 with the full control set ($p < 0.01$), closely matching the binary specifications. For hours, the maternal coefficient remains positive, ranging from 0.15 to 0.35 log points, but is imprecisely estimated under the continuous specification. The continuous labor results therefore reinforce the extensive-margin first stage more clearly than the hours response. The only exception is the provincial coefficient for hours in the full continuous specification, which is negative and significant. Because the same pattern does not appear for wage employment or in the binary exposure design, and because the continuous provincial measure captures broader local exposure rather than a specific maternal job shock, we interpret this estimate as specification-specific rather than as evidence against the maternal labor-supply channel.

For fertility, Table C9 confirms the negative maternal exposure effect. When maternal and paternal treatments are jointly included, the $\tau_m \times \text{Post}$ coefficient is -0.008 ($p < 0.05$), consistent with the binary estimate of -0.006 ($p < 0.01$). Controlling for paternal exposure sharpens the maternal estimate, producing coefficients between -0.008 and -0.012 that are statistically significant in columns (3)–(5). Paternal exposure ($\tau_d \times \text{Post}$) consistently raises fertility, with coefficients ranging from 0.021 to 0.026 ($p < 0.01$) across all combined specifications, consistent with the binary result that $F_i \times \text{Post}$ increases birth probability. The provincial continuous exposure also enters positively and significantly for fertility in columns (4)–(5), which is consistent with Prediction 3 if stronger local exposure captures broader income gains. We treat this result as suggestive rather than central because the provincial fertility response is not robust to the binary high-versus-low exposure contrast.

Across all three outcomes, the continuous specifications remain consistent with the model’s maternal channel rather than a generalized household income channel. These continuous-treatment estimates should nonetheless be interpreted with caution. As Callaway et al. (2024) show, the two-way fixed effects coefficient on a continuous dose is a weighted average of dose-specific average treatment effects, with implicit weights that can be negative for doses below the mean. We therefore do not interpret the continuous-treatment coefficient as a dose-response derivative. The binary specification delivers a well-defined high-versus-low exposure contrast under standard parallel trends, and remains our preferred estimand for both sign and magnitude. The continuous specifications are reported as a robustness check that the qualitative pattern survives when the full distribution of tariff-cut intensity is exploited.

7. Alternative Explanations

This section asks whether alternative explanations account for the full empirical pattern as well as the model in Section 2. A successful explanation must account for why maternal exposure raises male birth probability, increases maternal labor supply, and reduces fertility, while paternal and provincial exposure do not reproduce the sex-selection response. Table 1 summarizes the comparative statics that guide these tests. We begin with the pure income channel, which provides the clearest counterfactual because it predicts that paternal and broader local gains should also raise sex selection.

Under the income-effects explanation (Almond et al., 2019), any income source should increase sex selection, including paternal wages and broad local gains captured by provincial exposure. Across all specifications, paternal and provincial exposure coefficients for sex selection are small and statistically insignificant (Tables 3 and C7), even though these channels show some effects on other margins such as wage employment or fertility. The cross-channel pattern is difficult to reconcile with a model in which income alone drives sex selection.

The relative-returns mechanism (Qian, 2008; Rosenzweig and Schultz, 1982) predicts that expanding opportunities in female-intensive industries should reduce sex selection by making daughters more valuable. The maternal exposure coefficient is positive across all specifications (Tables 3 and C7), which is not consistent with this prediction. The result does not require denying that modern-sector opportunities raise the return to daughters' human capital. In our model, that force also raises desired investment in children, reinforcing the fertility-pressure logic emphasized by Jayachandran (2017) through the quantity-quality margin.

The preceding test interprets the relative-returns channel at the level of individual wages. In the model of Section 2, returns to daughters' human capital equal the maternal wage. Through the cost channel, higher w_m raises sex selection, while the value channel works in the opposite direction (Section 2.2). The net positive effect documented in our estimates suggests the cost channel dominates. But the relative-returns mechanism's original argument operates at a different level. It is not about whether a particular mother earns more today, but whether daughters born now will face better economic prospects as adults in the local labor market. When female-intensive industries expand in a province, all families in that province may expect daughters to find more rewarding work when they grow up, regardless of the mother's own industry. Appendix A.2.1 generalizes the model to allow province-level returns to daughters to rise independently of the mother's own wage. In that case, the local return shock can raise daughters' economic value without directly tightening the maternal time constraint. The predicted effect on sex selection is therefore more favorable to the relative-returns explanation than the individual maternal-wage test.

To test this local-market reading directly, we construct a female-weighted provincial BTA exposure measure $\tilde{P}_p^w = \sum_j f_j \cdot s_{jp} \cdot \tau_j$, where f_j is the pre-BTA female employment share in industry j , s_{jp} is the pre-BTA employment share of industry j in province p , and τ_j is the industry-level tariff

cut. The measure is high in provinces where female-intensive industries were initially important and received larger tariff cuts. We define $P_p^w = \mathbf{1}[\tilde{P}_p^w > \text{median}]$ as an indicator for above-median female-weighted provincial exposure. Under the relative-returns explanation, $P_p^w \times \text{Post}_t$ should reduce sex selection. We estimate:

$$y_{it} = \beta_1(M_i \times \text{Post}_t) + \beta_2(F_i \times \text{Post}_t) + \beta_3(P_{p(i)} \times \text{Post}_t) + \beta_4(P_{p(i)}^w \times \text{Post}_t) + \alpha_t + \delta_{p(i)} + \gamma_{j_{m(i)}} + \gamma_{j_{d(i)}} + X'_{it}\phi + \varepsilon_{it}, \quad (17)$$

The specification follows Equation (13) with $P_p^w \times \text{Post}_t$ added to the three baseline exposure channels. For fertility and labor outcomes, the fixed effects, controls, and post definitions follow the corresponding baseline specifications. The relative-returns prediction is $\beta_4 < 0$ for sex selection. Table C10 reports the results across all four outcomes. The female-weighted provincial coefficient is small and statistically insignificant for sex selection, hours, and fertility, with only a marginally significant 5 percentage point increase in maternal wage employment. The sign on sex selection is positive rather than negative. The local-market version of the relative-returns mechanism therefore does not account for the male-birth response.

Another potential explanation is bargaining. Unlike income or relative returns, bargaining does not make a sharp prediction from the direction of the wage shock alone because the effect depends on whose preferences receive more weight. A rise in maternal wages could increase mothers' bargaining power, and the resulting shift in household decision weights could affect fertility or sex selection even if childcare time costs did not change. To explain our main result, however, the bargaining channel would need stronger maternal bargaining power to raise son selection while also reducing fertility. That logic has two empirical implications. First, paternal wage gains should move bargaining power in the opposite direction and therefore should move sex selection, at least if the paternal shock is strong enough to affect household decisions. Second, households where mothers already have stronger bargaining positions should display more son selection before the trade shock.

The data do not support either implication. Paternal BTA exposure produces meaningful economic effects: fathers in high-exposure industries are 15 percentage points more likely to hold wage jobs (Table 4), and paternal exposure raises the annual birth probability by 0.5 percentage points (Tables 5 and C10). Yet the paternal exposure coefficient for sex selection is small and statistically insignificant across the binary (Table 3), continuous (Table C7), and panel (Table 6) specifications. We also test whether pre-BTA bargaining position predicts sex selection using the 1999 Census. Among births in 1997 and 1998, mothers with strictly more education than their husbands, our proxy for stronger maternal bargaining position, are no more likely to have a son (Table C12). These tests do not rule out all bargaining responses, but they make bargaining unlikely to be the main driver of the pattern we observe.

A final alternative is postnatal composition rather than prenatal sex selection. We observe children aged 0–1 who are alive and co-resident at the time of the survey, so maternal exposure

could in principle change the observed male share through sex-differential infant mortality, non-co-residence, or selective fostering. The mortality channel would require female infant mortality to rise relative to male infant mortality. Aggregate infant-mortality trends point in the opposite direction. Between 2002 and 2008, male and female infant mortality declined nearly symmetrically, and male mortality remained higher throughout (Appendix Table C13). Fostering of infants aged 0–1 is also rare in Vietnamese household structure. Administrative birth records linking ultrasound visits to maternal employment would provide a definitive test, but the available evidence suggests that postnatal composition is more likely a measurement limitation than the main explanation for the cross-channel pattern.

8. Concluding Remarks

This paper shows that better jobs for mothers can increase sex selection at birth when son preference is present and childcare remains primarily maternal. Higher maternal wages raise the opportunity cost of childbearing, reduce fertility, and increase pressure to concentrate births on sons. The quantity-quality response helps reconcile this pattern with rising investment in girls who are born: greater daughter returns can increase desired human-capital investment even as the sex ratio at birth becomes more male. Empirically, maternal exposure to the 2001 US-Vietnam Bilateral Trade Agreement raises male birth probability, increases maternal labor supply, and reduces annual birth hazards. Paternal and provincial exposure channels do not move sex selection, and panel estimates reduce concern that post-BTA sorting drives the result.

The mechanism we identify is not unique to Vietnam. Albania, Montenegro, and Georgia show similar trajectories around their own trade integration episodes (Figure 1), and the same structural conditions are present in many other developing economies with son preference, patrilocal family structures, and rapid expansion of female employment, including parts of Indonesia and sub-Saharan Africa. In Vietnam, the effects concentrate among older mothers and in provinces with higher pre-existing sex ratios, suggesting that the mechanism is most operative where son preference is strongest and mothers face a narrowing fertility horizon. More broadly, where mothers bear the primary childcare burden and cultural or institutional norms favor sons, new female employment opportunities can narrow economic gender gaps while widening demographic gender imbalance.

Addressing the link between female employment and sex selection requires policies that reduce the maternal time constraint directly. Income transfers or wage subsidies alone will not suffice, since the mechanism runs through the opportunity cost of mothers' time rather than household income. The human-capital extension shows that education costs can amplify the mechanism, but education-price policy is not a clean substitute for childcare because it can also change the relative cost of selecting sons. Future work with administrative birth records and direct measures of childcare access could test more precisely whether institutional childcare weakens the link between maternal employment opportunities and sex selection.

References

- ActionAid-Vietnam (2016). Unpaid Care Work: Make a House becomes a Home. Technical report, ActionAid Vietnam, Hanoi.
- Almond, D., Li, H., and Zhang, S. (2019). Land reform and sex selection in china. *Journal of Political Economy*, 127(2):560–585.
- Angrist, J. D. and Evans, W. N. (1998). Children and Their Parents' Labor Supply: Evidence from Exogenous Variation in Family Size. *American Economic Review*, 88(3):450–477.
- Anukriti, S. (2018). Financial incentives and the fertility-sex ratio trade-off. *American Economic Journal: Applied Economics*, 10(2):27–57.
- Anukriti, S., Bhalotra, S., and Tam, E. H. F. (2022). On the Quantity and Quality of Girls: Fertility, Parental Investments and Mortality. *The Economic Journal*, 132(641):1–36.
- Anukriti, S. and Kumler, T. J. (2019). Women's worth: Trade, female income, and fertility in India. *Economic Development and Cultural Change*, 67(3):687–724.
- Becker, G. S. and Lewis, H. G. (1973). On the Interaction between the Quantity and Quality of Children. *Journal of Political Economy*, 81(2, Part 2):S279–S288.
- Belanger, D. (2002). Son preference in a rural village in north Vietnam. *Studies in Family Planning*, 33(4):321–334.
- Borker, G., Eeckhout, J., Luke, N., Minz, S., Munshi, K., and Swaminathan, S. (2026). Wealth, Marriage, and Sex Selection. *American Economic Journal: Applied Economics*, 18(2):107–146.
- Callaway, B., Goodman-Bacon, A., and Sant'Anna, P. H. C. (2024). Difference-in-differences with a Continuous Treatment. *National Bureau of Economic Research*.
- Chakraborty, T. (2015). Trade Liberalization in a Traditional Society: Implications for Relative Female Survival. *World Development*, 74:158–170.
- Chao, F., Gerland, P., Cook, A. R., and Alkema, L. (2019). Systematic assessment of the sex ratio at birth for all countries and estimation of national imbalances and regional reference levels. *Proceedings of the National Academy of Sciences of the United States of America*, 116(19):9303–9311.
- Chen, Y., Li, H., and Meng, L. (2013). Prenatal sex selection and missing girls in China: Evidence from the diffusion of diagnostic ultrasound. *Journal of Human Resources*, 48(1):36–70.
- Doepke, M., Hannusch, A., Kindermann, F., and Tertilt, M. (2023). The economics of fertility: a new era. In *Handbook of the Economics of the Family*, volume 1, pages 151–254. Elsevier.

- Drozdoﬀ, N., McCaig, B., and Pavcnik, N. (2025). Foreign manufacturing jobs and women: Impacts on work, marriage, and childbirth. Working Paper.
- Duﬂo, E. (2003). Grandmothers and Granddaughters: Old-Age Pensions and Intrahousehold Allocation in South Africa. *World Bank Economic Review*, 17(1):1–25.
- Ebenstein, A. (2011). Estimating a Dynamic Model of Sex Selection in China. *Demography*, 48(2):783–811.
- General Statistics Office of Vietnam (2002–2010). Vietnam household living standards survey (vhlss), 2002–2010. Household survey microdata. Waves used in this paper: 2002, 2004, 2006, 2008, and 2010.
- Huynh, Q. and Ku, H. (2025). Gender Norms and Female Labor Supply: Evidence from Export Shocks in Vietnam. IZA Discussion Paper.
- Jayachandran, S. (2017). Fertility decline and missing women. *American Economic Journal: Applied Economics*, 9(1):118–139.
- Jayachandran, S. and Kuziemko, I. (2011). Why Do Mothers Breastfeed Girls Less than Boys? Evidence and Implications for Child Health in India. *The Quarterly Journal of Economics*, 126(3):1485–1538.
- Jayachandran, S. and Voena, A. (2026). Women’s power in the household. Forthcoming, Journal of Economic Literature. Draft dated January 14, 2026.
- Kovak, B. K. (2013). Regional eﬀects of trade reform: What is the correct measure of liberalization? *American Economic Review*, 103(5):1960–1976.
- Li, H., Yi, J., and Zhang, J. (2011). Estimating the Eﬀect of the One-Child Policy on the Sex Ratio Imbalance in China: Identification Based on the Diﬀerence-in-Diﬀerences. *Demography*, 48(4):1535–1557.
- McCaig, B. (2011). Exporting out of poverty: Provincial poverty in Vietnam and U.S. market access. *Journal of International Economics*, 85(1):102–113.
- McCaig, B., Nguyen, M., and Kaestner, R. (2022). Export expansion and investment in children’s human capital: Evidence from the u.s.-vietnam bilateral trade agreement. NBER Working Paper 29830, National Bureau of Economic Research.
- McCaig, B. and Pavcnik, N. (2018). Export markets and labor allocation in a low-income country. *American Economic Review*, 108(7):1899–1941.
- McCaig, B., Pavcnik, N., and Wong, W. F. (2025). Foreign and Domestic Firms: Long-Run Employment Eﬀects of Trade Liberalization. *Review of Economics and Statistics*. Forthcoming.

- Mitra, D., Pham, H., and Ural Marchand, B. (2026). Exports and intergenerational mobility. *Journal of International Economics*, 161:104239.
- Qian, N. (2008). Missing women and the price of tea in China: The effect of sex-specific earnings on sex imbalance. *Quarterly Journal of Economics*, 123(3):1251–1285.
- Rosenzweig, M. R. and Schultz, T. P. (1982). Market opportunities, genetic endowments, and intrafamily resource distribution: child survival in rural India. *American Economic Review*, 72(4):803–815.
- Rosenzweig, M. R. and Wolpin, K. I. (1980). Testing the Quantity-Quality Fertility Model: The Use of Twins as a Natural Experiment. *Econometrica*, 48(1):227.
- Rosenzweig, M. R. and Zhang, J. (2009). Do population control policies induce more human capital investment? Twins, birth weight and China's "one-child" policy. *Review of Economic Studies*, 76(3):1149–1174.
- Ruggles, S., Cleveland, L., Lovaton, R., Sarkar, S., Sobek, M., Burk, D., Ehrlich, D., Heimann, Q., Lee, J., and Merrill, N. (2025). Integrated public use microdata series, international: Version 7.6. Dataset.
- Sobek, M. and Kennedy, S. (2009). The Development of Family Interrelationship Variables for International Census Data. *Minnesota Population Center Working Paper Series*.
- STAR-Vietnam (2003). An Assessment of the Economic Impact of the United States - Vietnam Bilateral Trade Agreement. Technical report, The National Political Publishing House, Hanoi.

Appendix to Better Jobs for Mothers, More Sons at Birth

Nguyễn Q. Huỳnh and Ngọc T. T. Nguyen

A. Model Appendix

This appendix provides the full derivations supporting the model in Section 2. Appendix A.1 derives the first-order conditions and comparative statics for the baseline model via Slutsky decomposition. Appendix A.2 extends the model to include human capital investment in a quantity-quality framework.

A.1 Baseline Model Derivations

A.1.1 Lagrangian and First-Order Conditions

The household maximizes $U(c, \ell, n, B)$ with $B = bn$, choosing (c, ℓ, n, b) subject to the budget constraint $c + w_m \ell + w_m \gamma n + p_b b n = W$. The Lagrangian is

$$\mathcal{L} = U(c, \ell, n, bn) + \lambda [W - c - w_m \ell - w_m \gamma n - p_b b n]. \quad (\text{A1})$$

The first-order conditions with respect to c , ℓ , n , and b are

$$U_c = \lambda, \quad (\text{A2})$$

$$U_\ell = \lambda w_m, \quad (\text{A3})$$

$$U_n + U_B b = \lambda (w_m \gamma + p_b b), \quad (\text{A4})$$

$$U_B = \lambda p_b. \quad (\text{A5})$$

The fertility condition (A4) includes $U_B b$ on the left (the marginal benefit of one more child through the boy channel, since each child yields b expected boys) and $p_b b$ on the right (the sex-selection cost per child). Substituting Equation (A5) into Equation (A4) yields

$$U_n = \lambda w_m \gamma. \quad (\text{A6})$$

At the optimum, the effective shadow price of fertility is $w_m \gamma$, the maternal time cost alone. The household simultaneously optimizes b , so the sex-selection cost of an additional child ($p_b b$) is exactly offset by the marginal boy value ($U_B b$). Dividing through by $U_c = \lambda$, the MRS conditions reported in the main text (Equation 4) follow.

The effective price structure $(1, w_m, w_m \gamma, p_b)$ for goods (c, ℓ, n, B) permits standard Slutsky decomposition.

Maternal wage effect on sex selection. The maternal wage w_m enters two prices, $p_\ell = w_m$ and $p_n = w_m\gamma$. Thus, the compensated effect on the sex ratio $b = B/n$ operates through both channels:

$$\frac{\partial b^S}{\partial w_m} = \underbrace{1 \cdot \frac{\partial b^S}{\partial p_\ell}}_{\text{leisure-price channel}} + \underbrace{\gamma \cdot \frac{\partial b^S}{\partial p_n}}_{\text{fertility-price channel}} > 0, \quad (\text{A7})$$

where positivity holds when the boy share b and fertility n are Hicksian substitutes ($\partial b^S/\partial p_n > 0$). The household raises the boy share rather than just cutting the number of children when children become more expensive. The same logic appears in the literature (Jayachandran, 2017): with son preference, smaller families concentrate desired sons into a higher share. The total (Marshallian) effect adds the income effect:

$$\frac{\partial b}{\partial w_m} = \underbrace{\frac{\partial b^S}{\partial w_m}}_{> 0 \text{ (substitution)}} + \underbrace{h_m \frac{\partial b}{\partial W}}_{\text{income effect on } b}, \quad (\text{A8})$$

where $h_m = 1 - \ell - \gamma n$ is the mother's labor supply. The income effect on b is equivalent to $h_m \left(\frac{1}{n} \frac{\partial B}{\partial W} - \frac{b}{n} \frac{\partial n}{\partial W} \right)$ via the chain rule on $b = B/n$. In elasticity terms, $\eta_b^W = \eta_B^W - \eta_n^W$: the income effect on b is positive when $\eta_B^W > \eta_n^W$ (boys more income-elastic than fertility), zero when they coincide, and negative otherwise. Under son preference, it is likely that marginal dollars tilt toward the boy share, so $\eta_B^W \geq \eta_n^W$. With substitution positive and income weakly positive, $\partial b/\partial w_m > 0$.

Maternal wage effect on fertility. Both price channels reduce fertility:

$$\frac{\partial n}{\partial w_m} = \underbrace{\frac{\partial n^S}{\partial w_m}}_{< 0 \text{ (substitution)}} + \underbrace{h_m \frac{\partial n}{\partial W}}_{\text{income effect}}. \quad (\text{A9})$$

The substitution effect is unambiguously negative: own-price Slutsky $\partial n^S/\partial p_n \leq 0$ enters the chain rule with coefficient γ . The income effect $h_m \cdot \partial n/\partial W$ may partially offset if n is a normal good, but its force is scaled by h_m , the mother's net labor supply. When γ is significant because childcare is materially time-intensive, the substitution term dominates and $\partial n/\partial w_m < 0$.

Maternal wage effect on labor supply. The time constraint $h_m = 1 - \ell - \gamma n$ pins down labor supply as a residual. Substituting the Slutsky decompositions of $\partial \ell/\partial w_m$ and $\partial n/\partial w_m$:

$$\frac{\partial h_m}{\partial w_m} = - \underbrace{\left[\frac{\partial \ell^S}{\partial w_m} + \gamma \frac{\partial n^S}{\partial w_m} \right]}_{\text{substitution: } > 0} - \underbrace{h_m \left[\frac{\partial \ell}{\partial W} + \gamma \frac{\partial n}{\partial W} \right]}_{\text{income: } < 0 \text{ if } \ell, n \text{ normal}}. \quad (\text{A10})$$

This decomposition exhibits the classic backward-bending labor supply tension where substitution and income work in opposite directions. The same condition that signs $\partial n / \partial w_m$, i.e. substitution dominates because γ amplifies the time-price terms, signs labor supply: $\partial h_m / \partial w_m > 0$.

Paternal wage effect. The paternal wage w_d enters no price. Its effect on any demand operates entirely through full income W :

$$\frac{\partial b}{\partial w_d} = \frac{\partial b}{\partial W}. \quad (\text{A11})$$

By the same elasticity decomposition $\eta_b^W = \eta_B^W - \eta_n^W$ used for the maternal income effect, $\partial b / \partial W$ is positive when boys are more income-elastic than fertility, zero when the elasticities coincide, and negative otherwise. Under son preference, the positive case is plausible because marginal dollars may tilt toward birth composition. The critical asymmetry does not require this income effect to be large. The maternal wage adds substitution effects on top of the income effect, while the paternal wage has income only.

The paternal wage, via income effect, also raises fertility ($\partial n / \partial w_d > 0$ when n is normal) and reduces maternal hours ($\partial h_m / \partial w_d < 0$, since leisure is normal). Non-labor income I has the identical structure where $\partial b / \partial I = \partial b / \partial W$ and likewise for the other outcomes.

A.2 Human Capital Extension

We now extend the baseline model to let parents also choose how much to invest in each child, separately for sons and daughters. The choice set expands from (c, ℓ, n, b) to $(c, \ell, n, b, e_b, e_g)$, where e_b and e_g are the educational inputs (e.g., years of schooling) per son and per daughter, purchased at unit price p_e . Each unit of input is converted into human capital through a production function $h(\cdot)$ with $h' > 0$, $h'' < 0$, and $h(0) = 0$. The return to a unit of human capital is r_b for sons and r_g for daughters; in the baseline mapping, $r_b = w_d$ and $r_g = w_m$. Average per-child human capital is $\bar{E} = b r_b h(e_b) + (1 - b) r_g h(e_g)$, and parents value it through a separable, increasing, concave function $V(\bar{E})$.

The household problem is to maximize $U(c, \ell, n, B) + V(\bar{E})$, with $B = bn$, subject to the budget constraint

$$c + w_m \ell + n[w_m \gamma + p_e(b e_b + (1 - b) e_g)] + p_b b n = W. \quad (\text{A12})$$

Relative to the baseline budget, education spending $p_e[b e_b + (1 - b) e_g]$ enters per child, scaling with n . The Lagrangian is

$$\mathcal{L} = U(c, \ell, n, B) + V(\bar{E}) - \lambda [c + w_m \ell + n(w_m \gamma + p_e(b e_b + (1 - b) e_g)) + p_b b n - W]. \quad (\text{A13})$$

The first-order conditions for education investment are

$$\frac{\partial \mathcal{L}}{\partial e_b} = V'(\bar{E}) \cdot b r_b h'(e_b) - \lambda n b p_e = 0, \quad (\text{A14})$$

$$\frac{\partial \mathcal{L}}{\partial e_g} = V'(\bar{E}) \cdot (1 - b) r_g h'(e_g) - \lambda n (1 - b) p_e = 0. \quad (\text{A15})$$

Dividing by b and $(1 - b)$ respectively and equating yields the education condition stated in the main text:

$$r_b h'(e_b^*) = r_g h'(e_g^*). \quad (\text{A16})$$

Because $h'' < 0$, the gender with higher returns receives more investment. When men earn more than women in the labor market ($w_d > w_m$), the marginal return to a son's human capital exceeds the marginal return to a daughter's, so optimal investment per son exceeds optimal investment per daughter ($e_b^* > e_g^*$). Before 2005, the female-to-male ratio of university enrollment lay well below one (Figure 1, panel d), consistent with parents allocating more education to sons than to daughters under the prevailing gender wage gap.

The first-order condition for b is

$$U_B \cdot n + V'(\bar{E}) \cdot [r_b h(e_b) - r_g h(e_g)] = \lambda n [p_b + p_e(e_b - e_g)]. \quad (\text{A17})$$

The left-hand side is the marginal benefit of raising b . The first term ($U_B \cdot n$) is the direct preference for sons. The second term captures how shifting the sex composition toward boys changes average human capital. When sons are more valuable than daughters ($r_b h(e_b) > r_g h(e_g)$), this term is positive, adding to the incentive for sex selection. The right-hand side is the marginal cost: the sex-selection technology cost p_b plus the education premium $p_e(e_b - e_g)$ for each son over each daughter, summed over n children.

The first-order condition for n is

$$U_n + U_B b = \lambda [w_m \gamma + p_e(b e_b + (1 - b) e_g) + p_b b]. \quad (\text{A18})$$

Relative to the baseline fertility condition (Equation A4), the right-hand side now includes the per-child education cost $p_e(b e_b + (1 - b) e_g)$. An additional child requires not only the time cost $w_m \gamma$ and the sex-selection cost $p_b b$ but also investment in the new child's human capital, weighted by the expected sex composition. This term is the channel through which education prices feed back into the cost of fertility.

From Equation (A14), dividing by b :

$$V'(\bar{E}) \cdot r_b h'(e_b) = \lambda n p_e. \quad (\text{A19})$$

From Equation (A15), dividing by $(1 - b)$:

$$V'(\bar{E}) \cdot r_g h'(e_g) = \lambda n p_e. \quad (\text{A20})$$

These expressions show that sex composition does not enter the education FOCs directly once

we scale by the relevant child share. It enters indirectly through the shadow value λ . When $r_b > r_g$, optimal investment satisfies $e_b^* > e_g^*$. Holding (n, e_b^*, e_g^*) fixed, the education bill is $C_E = np_e[be_b^* + (1-b)e_g^*]$, so

$$\frac{\partial C_E}{\partial b} = np_e(e_b^* - e_g^*) > 0.$$

Replacing a daughter with a son raises the education bill when $e_b^* > e_g^*$, so a son carries a quality premium beyond the sex-selection cost p_b . Symmetrically, higher investment per child raises the marginal cost of an additional child through (A18). This two-way dependence is the quantity-quality tradeoff of [Becker and Lewis \(1973\)](#), and it is the channel through which the education margin enters both the fertility price and the sex-selection price below.

Rewriting the budget as in Equation (10), the effective prices are

$$p_c = 1, \quad p_\ell = w_m, \quad p_n = w_m\gamma + p_e e_g, \quad p_B = p_b + p_e(e_b - e_g). \quad (\text{A21})$$

Relative to the baseline (where $p_n = w_m\gamma$ and $p_B = p_b$), the maternal wage now enters p_n through two channels and the paternal wage enters p_B through the son education cost differential. The comparative statics follow the Slutsky decomposition of [Appendix A.1](#), applied to these modified prices.

Maternal wage effect on sex selection. A rise in w_m raises $r_g = w_m$. Holding the shadow cost of quality fixed, the direct return effect raises daughter investment because $h'' < 0$, narrowing the education gap. Indirect feedback through the shadow value of income λ and fertility n also moves both e_b^* and e_g^* . Provided the direct return effect dominates this feedback, the effect of w_m on prices is

$$\frac{dp_n}{dw_m} = \gamma + p_e \frac{de_g^*}{dw_m} > \gamma, \quad (\text{A22})$$

$$\frac{dp_B}{dw_m} = p_e \left(\frac{de_b^*}{dw_m} - \frac{de_g^*}{dw_m} \right). \quad (\text{A23})$$

The fertility price p_n then rises by a larger magnitude than in the baseline (where $dp_n/dw_m = \gamma$), because daughter education costs increase alongside time costs. The sex-selection price p_B declines when the education gap $e_b^* - e_g^*$ narrows.

In the baseline model, w_m enters only through prices in the budget constraint, and the Slutsky decomposition fully captures the substitution effect. With the human capital extension, w_m also enters the utility function through $r_g = w_m$ in $V(\bar{E})$. The maternal wage therefore affects b through two distinct routes. The first route operates through the budget constraint: w_m changes the effective prices, and the household reoptimizes along the budget. The price-side effect paral-

lels Equation (A7) with the modified prices:

$$\left. \frac{\partial b^S}{\partial w_m} \right|_{\text{HC}} = \underbrace{1 \cdot \frac{\partial b^S}{\partial p_\ell}}_{\text{leisure price}} + \underbrace{\frac{dp_n}{dw_m} \cdot \frac{\partial b^S}{\partial p_n}}_{\text{fertility price (amplified)}} + \underbrace{\frac{dp_B}{dw_m} \cdot \frac{\partial b^S}{\partial p_B}}_{\text{sex-selection price (new)}}. \quad (\text{A24})$$

In the baseline model, only the first two terms appear. The human capital extension adds the education component to the fertility-price channel and introduces the third term. We call this entire price-side mechanism the *cost channel*. Rising w_m makes children more expensive to raise and sons relatively cheaper, pushing b up. The first two terms of the cost channel are positive under the baseline condition that fertility and boys are Hicksian substitutes ($\partial b^S / \partial p_n > 0$). The third term is positive when the compensated own-price effect on boys is negative ($\partial b^S / \partial p_B < 0$) and the education gap narrows ($dp_B / dw_m < 0$). When both conditions hold, all three cost-channel forces reinforce.

The second route operates through the utility function, not the budget. Because $r_g = w_m$, the maternal wage also enters $V(\bar{E})$ directly. From the b -FOC (Equation A17), the marginal benefit of sex selection includes $V'(\bar{E}) \cdot [r_b h(e_b) - r_g h(e_g)]$. When w_m rises, r_g increases, so daughters' human capital $r_g h(e_g^*)$ rises and the gap $r_b h(e_b) - r_g h(e_g)$ narrows. The marginal benefit of selecting for sons falls. We call this the *value channel*. When daughters' future earnings prospects improve, each daughter becomes more valuable, and the household has less reason to skew births toward sons.

The total effect of w_m on b combines both routes:

$$\frac{db}{dw_m} = \underbrace{\left. \frac{\partial b^S}{\partial w_m} \right|_{\text{HC}}}_{\substack{\text{cost channel} \\ (+, \text{ through prices})}} + \underbrace{\left. \frac{\partial b}{\partial r_g} \right|_{\text{prices fixed}}}_{\substack{\text{value channel} \\ (-, \text{ through } V)}} + \text{income effect}. \quad (\text{A25})$$

These forces deliver the positive maternal prediction in Table 1. Higher maternal wages raise the time price of children, raise the education component of the fertility price, and lower the relative education-cost penalty of selecting sons when the gender investment gap narrows. The daughter-value channel works in the opposite direction, but it is unlikely to overturn these price-side forces to make the maternal effect zero or negative. The model therefore predicts $\partial b / \partial w_m > 0$ in settings where childcare remains materially maternal and fertility can adjust.

Maternal wage effect on fertility. In the baseline, $dn / dw_m < 0$ when the substitution effect dominates the income effect. The human capital extension strengthens this substitution force when daughter investment rises with the maternal wage, because the fertility price rises from $p_n = w_m \gamma$ to $p_n = w_m \gamma + p_e e_g$.

Paternal wage effect. As with the maternal wage, w_d affects b through both prices and utility. On the price side, w_d enters p_B directly through sons' returns and can enter p_n indirectly through the shadow-cost feedback:

$$\frac{dp_B}{dw_d} = p_e \left(\frac{de_b^*}{dw_d} - \frac{de_g^*}{dw_d} \right), \quad \frac{dp_n}{dw_d} = p_e \frac{de_g^*}{dw_d}. \quad (\text{A26})$$

Higher w_d raises sons' returns ($r_b = w_d$), so it directly raises e_b^* and widens the gap $e_b - e_g$, raising the sex-selection price p_B . The fertility price p_n moves only through indirect feedback on e_g^* and is therefore second-order. The price-side channel thus has a primary negative force through p_B and a possible offsetting positive force through p_n :

$$\left. \frac{\partial b^S}{\partial w_d} \right|_{\text{HC}} = \underbrace{\frac{dp_B}{dw_d} \cdot \frac{\partial b^S}{\partial p_B}}_{\text{sons costlier}} + \underbrace{\frac{dp_n}{dw_d} \cdot \frac{\partial b^S}{\partial p_n}}_{\text{quality-price feedback}}. \quad (\text{A27})$$

The first term is negative when sons become more expensive to educate. The second term, operating through the indirect quality-price channel, is positive when p_n rises. On the utility side, w_d enters $V(\bar{E})$ through $r_b = w_d$. When w_d rises, sons' human capital $r_b h(e_b^*)$ increases, widening the gap $r_b h(e_b) - r_g h(e_g)$ in the b -FOC (Equation A17). The marginal benefit of selecting for sons rises. This value channel is positive, the mirror image of the maternal case. The income effect is also weakly positive (Equation A11). The total effect is therefore:

$$\frac{db}{dw_d} = \underbrace{\left. \frac{\partial b^S}{\partial w_d} \right|_{\text{HC}}}_{\substack{\text{cost channel} \\ (-)}} + \underbrace{\left. \frac{\partial b}{\partial r_b} \right|_{\text{prices fixed}}}_{\substack{\text{value channel} \\ (+)}} + \text{income effect (+)}. \quad (\text{A28})$$

The human-capital extension therefore makes the paternal prediction weaker than the baseline pure-income prediction. Higher paternal wages raise sons' returns, which increases the value of sons but also raises the education-cost premium of selecting sons. The value channel and the income effect push b up, while the son-cost channel pushes b down. Because higher paternal wages do not raise the maternal time price of children, they lack the main force that drives the maternal prediction. The model therefore predicts a small or null paternal effect on sex selection.

The fertility prediction $\partial n / \partial w_d > 0$ is likely positive. The direct income effect raises fertility, as in the baseline. The human capital extension can introduce a partially offsetting force if higher w_d raises p_n through the indirect quality-price channel (Equation A26), creating a substitution penalty on quantity. Because the direct positive income effect is the leading force and the indirect quality-price feedback is second-order, the net effect on fertility is likely positive.

Non-labor income effect. Non-labor income I enters only the budget constraint. The cost and value channels that operate through parental wages are absent for I , leaving only the income

effect: $\partial b/\partial I = \partial b/\partial W$ and $\partial n/\partial I = \partial n/\partial W$. Under the income-elasticity condition above, the effect on b is weakly positive. Without any price-side force, the model predicts a small effect on sex selection and a positive effect on fertility.

A.2.1 When Returns Differ from Individual Wages

The main text sets $r_g = w_m$ and $r_b = w_d$, identifying children's returns with same-sex parental wages. This subsection relaxes that restriction. When female-intensive industries expand in a province, the local return to daughters' human capital may rise for all families, regardless of individual mothers' industries. Let r_g now denote this province-level return, which may differ from any individual mother's wage w_m .

The education FOCs carry through with r_g and r_b as separate parameters. The education condition becomes $r_b h'(e_b^*) = r_g h'(e_g^*)$, and the quantity-quality interaction (Equation A20) still determines optimal investment, so e_g^* depends on both r_g and n . The modified prices become $p_n = w_m \gamma + p_e e_g(r_g, n)$ and $p_B = p_b + p_e(e_b(r_b, n) - e_g(r_g, n))$. The key difference from the individual-level case is that w_m still determines the time cost of childcare ($w_m \gamma$ in p_n), but r_g now independently determines the education investment in daughters.

When r_g rises independently of w_m , two opposing forces emerge. The cost channel raises sex selection: higher r_g increases optimal investment e_g^* , raising the fertility price p_n and lowering the sex-selection price p_B . The value channel lowers sex selection: daughters' economic value rises in $\bar{E}$, reducing the household's willingness to select for sons (through the b -FOC, Equation A17). Unlike the maternal wage case, there is no time-cost force ($w_m \gamma$) giving the cost channel a built-in advantage, because r_g moves independently of w_m . The net effect is therefore more ambiguous:

$$\frac{db}{dr_g} = \underbrace{\frac{\partial b}{\partial r_g} \Big|_{\text{cost}}}_{>0} + \underbrace{\frac{\partial b}{\partial r_g} \Big|_{\text{value}}}_{<0}. \quad (\text{A29})$$

This decomposition motivates the empirical test using female-weighted provincial exposure P_p^w in Section 7. The measure P_p^w captures province-level variation in returns to female labor that operates independently of individual maternal wages. If the cost and value channels approximately offset each other, P_p^w can improve female labor-market opportunities, and potentially raise desired investment in daughters through e_g^* , without producing a detectable net effect on sex selection. The maternal wage effect can still be positive because, when $r_g = w_m$, the shock also raises the maternal time cost of children; P_p^w shifts daughters' local returns without directly tightening that maternal time constraint.

A.2.2 Binding Fertility Constraints

The baseline model treats fertility as a choice variable. When policy constrains fertility, as under China's one-child policy, the fertility margin is restricted and the model's predictions change. Suppose n is fixed at $\bar{n}$. The household chooses (c, ℓ, b, e_b, e_g) to maximize $U(c, \ell, \bar{n}, B) + V(\bar{E})$ subject to

$$c + w_m \ell + \bar{n} [w_m \gamma + p_e (b e_b + (1 - b) e_g)] + p_b b \bar{n} = W. \quad (\text{A30})$$

The education FOCs are unchanged. The education ratio (Equation A16) and the quantity-quality interaction (Equation A20) both carry through with $\bar{n}$ replacing n . The critical difference is that $\bar{n}$ is a parameter, not a choice variable. From Equation (A20), the direct shadow cost of quality is $\lambda \bar{n} p_e$. Holding the shadow value of income fixed, a lower $\bar{n}$ reduces this direct cost and supports higher investment per child. When the constraint binds at a low level (such as one child), this quantity-quality force pushes toward heavier investment in each child. If sons have higher returns ($r_b > r_g$), the education FOC implies $e_b^* > e_g^*$ and raises the education-cost premium for selecting a son, $p_e(e_b^* - e_g^*)$.

The constrained model generates three predictions that differ from the unconstrained case. First, constrained fertility raises the level incentive for sex selection but changes the composition of the maternal wage effect. The b -FOC under fixed fertility is

$$U_B + \frac{V'(\bar{E})}{\bar{n}} [r_b h(e_b) - r_g h(e_g)] = \lambda [p_b + p_e (e_b - e_g)]. \quad (\text{A31})$$

A low $\bar{n}$ raises the marginal value of an additional boy when preferences over boys are concave, because $B = b\bar{n}$ is small and $U_B(B)$ is high. It also scales up the human-capital composition term by $1/\bar{n}$. These forces make sex selection more intense in levels under a tight fertility cap.

The marginal effect of a maternal wage increase is different. In the unconstrained model, the fertility-price channel ($\gamma \cdot \partial b^S / \partial p_n$) is a primary positive force driving b up when w_m rises: higher maternal wages make another child more costly, fertility falls, and son preference is concentrated over fewer births. When n is fixed at $\bar{n}$, the time cost $w_m \gamma \bar{n}$ becomes a fixed expenditure for the constrained number of births. Since households cannot reduce n in response to that higher cost, the fertility-price channel is shut off. The remaining forces are the leisure-price channel, the sex-selection price channel (from changes in the education gap), the income effect, and the value channel from daughters' returns. Thus the constrained model can have a higher level of sex selection even though the marginal maternal-wage effect is not necessarily larger than in the unconstrained model.

Second, the relative-returns channel changes as well. When province-level returns r_g rise independently of w_m (as in the preceding subsection), the cost and value channels still oppose each other (Equation A29), but the constrained setting changes their relative force. Because n is fixed, higher daughter investment no longer raises the price of an additional child; the remaining cost-

side effect operates mainly through the sex-selection price p_B . At the same time, investment per child is high, so the return to a daughter $r_g h(e_g^*)$ is large. A marginal increase in r_g therefore produces a larger increase in daughters' economic value, amplifying the value channel and making a negative effect of r_g on sex selection more likely.

Third, the income effect on sex selection is likely positive. In the unconstrained model, non-labor income I has a negligible effect on b because it does not enter any effective price. Under constrained fertility, the income gain that would have raised n cannot do so. If sons are income-elastic at the margin ($\eta_B^W > \eta_n^W$), part of the blocked fertility response may instead appear as higher b .

This constrained-fertility case also clarifies how to read the China evidence. With a low fertility cap, parents cannot respond to son preference by having more children until a son is born. The remaining adjustment margin is the sex composition of the permitted births. This is why evidence on China's one-child policy links tighter fertility restrictions to higher sex ratios (Li et al., 2011; Ebenstein, 2011). The same logic makes income shocks more likely to affect sex selection: when additional income cannot raise fertility, it can instead raise the demand for sons, consistent with Almond et al. (2019). Qian's result points in the opposite direction because the shock is different. Higher tea prices increased returns in a female-intensive agricultural activity, raising girls' relative economic value rather than the maternal time cost of childrearing.¹⁷ Under constrained fertility, that value channel can be especially powerful because each child receives more parental attention and investment. The contrast with Vietnam is therefore not a contradiction: Vietnam has an active fertility margin and a maternal time-cost shock, while the China results come from settings where fertility is constrained or where the shock changes girls' relative economic value more directly.

A.2.3 Policy Implications: Childcare versus Education Prices

The effective prices clarify why the model yields a cleaner policy implication for childcare than for education prices. The maternal childcare parameter γ enters only the fertility price, $p_n = w_m \gamma + p_e e_g$. Reducing γ therefore directly weakens the channel through which higher maternal wages make each child more costly, reduce fertility, and intensify sex selection.

The education price p_e is less direct because it affects two prices at once. It enters the fertility price p_n , so a higher education price raises the cost of an additional child. It also enters the sex-selection price, $p_B = p_b + p_e(e_b - e_g)$, because replacing a daughter with a son changes the education bill when boys and girls receive different investment. When $e_b > e_g$, lowering p_e makes children cheaper but also lowers the education-cost penalty for selecting a son. Optimal investments e_b^* and e_g^* also adjust when p_e changes. The model therefore does not imply that lowering education prices is an equivalent substitute for reducing maternal childcare time. Education costs

¹⁷Qian (2008) studies sex-specific agricultural income shocks in rural China, using tea as a female-intensive crop and orchards as a male-intensive crop. The mechanism is therefore not an education-return shock in the narrow sense, nor does it require that mothers' wages rise one-for-one with tea prices. In the notation here, it is best interpreted as a shock to the relative economic value of girls, holding the maternal childcare time cost conceptually distinct.

help explain why maternal wage gains can amplify the fertility response, but the cleaner policy lever in this framework is the maternal time cost of children.

B. Data Appendix: Family Interrelationship Algorithm

Variables categorizing family interrelationships within the household are crucial in this paper. In particular, we need “pointer” variables identifying each person’s mother, father, and spouse. VHLSS has a variable that classifies each individual’s relationship to the household head. However, this variable is ambiguous and erroneous for the following reasons. First, extended family is common in Vietnam: roughly 34.91% of the VHLSS sample lives in the same household with relatives other than their nuclear family members. Because VHLSS does not identify subfamilies within the extended household, the “relate-to-head” variable will miss the interrelationship of subfamilies other than the head’s. Second, with the exception of VHLSS 2002, all other waves do not identify “children-in-law.” An immediate consequence is missing spousal links and therefore parental links in many cases. Thus, we write an algorithm to improve the precision of parent-child and spousal pairings in VHLSS.

We adapt a similar algorithm that has been used by the Minnesota Population Center ([Sobek and Kennedy, 2009](#)) to generate “pointer” variables for IPUMS–International Census Data. We utilize the four variables: relationship to household head, age, and marital status, in combination with the relative position of household members in the roster listing, to infer relationships. We first establish spousal links (generate variable SPLOC); then we find the mother and father for each individual (generate variables MOMLOC and POPLOC). Regarding parental linkage, we look for the mother before looking for the father. Once we find the mother for a child, we assign her husband as the father of this child. Only when we cannot find a child’s mother do we locate the child’s father and assign the father’s wife as the mother.

Tables [B1](#) and [B2](#) document the rules we apply to these matching tasks. Our algorithm applies the rules sequentially. If the first rule finds a match for a given individual, the second rule no longer applies to that person, and so on. Whenever there are ambiguous multiple potential spouses or multiple potential parents, we drop the entire household from the sample.

Fortunately, VHLSS 2014 and 2016 provide variables that identify the biological fathers and mothers of children under 16. We use this information to validate our algorithm and report the results in Table [B3](#). The algorithm identifies both biological parents correctly for 88.95% of children. Looking at each parent separately, it identifies the correct mother for 93.23% of children and the correct father for 90.99%, allowing the other parent link to be incorrect. The joint match rate is mechanically lower than the single-parent match rates because both parent links must be correct at the same time. These validation results suggest that the algorithm recovers parent-child and spousal links with sufficient accuracy for the empirical analysis.

Table B1: Rules for SPLOC Construction

Rule	Individual's relationship to head	Partner's relationship to head	Age difference	Both Married	Require adjacency	Only applicable to 2002	Notes
Strong couple pairing, couple adjacency preferred							
	Head	Spouse	No	No	No		1
	Parent	Parent	No	Yes	Yes		1
	Grandparent	Grandparent	No	Yes	Yes		1
	Child	Child-in-law	No	Yes	Yes	Yes	1
Weak couple pairing, couple adjacent							
	Grandchild	Grandchild	Yes	Yes	Yes		1, 2
	Other relationship	Other relationship	Yes	Yes	Yes		1, 2
	Sibling	Sibling	Yes	Yes	Yes	Yes	1, 2
	Grandchild	Other relationship	Yes	Yes	Yes		1, 2
	Sibling	Other relationship	Yes	Yes	Yes	Yes	1, 2
	Child	Other relationship	Yes	Yes	Yes		1, 2
Weak couple pairing, special type child-child							
	Child	Child	Yes	Yes	Yes		1, 2, 6
Weak couple pairing, couple not adjacent							
	Child	Child-in-law	No	Yes	Closest proximity	Yes	5
	Child	Other relationship	Yes	Yes	No		2, 3
	Child	Child	Yes	Yes	No		1, 2, 4, 6

Notes:

1. Drop the entire household if there is any person that could be assigned to 2 couples by the adjacency rule.
2. A woman can be no more than 20 years older or 35 years younger than a potential male partner.
3. For non-adjacent couple pairing, among the potential spouses who satisfy the age and marital requirements, select the person who is closest in age and impose that the husband is older than the wife.
4. For child-child non-adjacent couple pairing, drop all households where this rule yields multiple potential spouses.
5. These non-adjacent couples are matched based on having closest proximity to each other.
6. Once the couple is identified, the in-law is distinguished. For child-child couples, assume the first listed spouse is the biological child to the household head, hence the second listed spouse is the child-in-law.

C. Additional Tables and Figures

Table B2: Rules for Child-Parent Relationship Construction (after SPLOC is Generated)


Rule	Child's relationship to head	Parent's relationship to head	Age difference	Proximity requirement applicable	Only applicable to 2002	Notes
Links involving Head, Spouse, and Grandparent (unambiguous)						
	Child	Head, spouse	No	No		
	Child-in-law	Head, spouse	No	No	Yes	4
	Head	Parent	No	No		
	Spouse	Parent	No	No		
	Sibling	Parent	No	No	Yes	
	Parent	Grandparent	No	No		
Links between grandchildren and children						
	Grandchild	Child, child-in-law	15–44	Weak		1
	Grandchild	Other relationship	15–44	Weak		1, 2
Links involving other relatives						
	Other relationship	Grandchild	15–44	Weak		1
	Other relationship	Other relative	15–44	Weak		1
	Other relationship	Sibling	15–44	Weak		1
	Other relationship	Child, child-in-law	15–44	Weak		1, 3
Notes:						
1. Weak proximity requires that the child must be listed after its potential mothers (or potential fathers if it has no potential mothers). Among them, its mother is the one listed closest to the child.						
2. Impose that no person with code "Child" is present in the household. The mother of the grandchildren in these cases tend to be listed as "Others" since there is no category for "Child-in-law."						
3. Impose that no person with code "Grandchild" is present in the household. These cases tend to mix up the numerical code for "Grandchild" (code 6) and that for "Other relationship" (code 7).						
4. The head and spouse are the child-in-law's parents-in-law, not biological parents. This link identifies the household structure (the child-in-law married into the head's family) rather than a biological parent-child relationship. Children-in-law who have their own infants in the household are matched to those infants through the grandchild rules below.						

Table B3: Performance of our algorithm in locating parents

	Freq.	Percent	Cum.
Both mom and dad correct	80021	88.95	88.95
Dad correct, mom incorrect	1831	2.04	90.99
Dad incorrect, mom correct	3853	4.28	95.27
Both mom and dad incorrect	4253	4.73	100.00
Total	89958	100.00	


Notes: This table compares the parent links generated by our algorithm with the biological parent links provided by VHLSS 2014–2016 for children under 16.

Figure C1: Correlations between Treatment Variables by Survey Wave




Notes: Pairwise correlations between continuous tariff-cut exposure measures by survey wave. Mother (τ_m) is the mother's industry tariff cut, Father (τ_d) is the father's industry tariff cut, and Province (τ_p) is shift-share provincial exposure. Data are from the 1999 Census baseline and VHLSS 2002–2010.

Figure C2: Individual Event Study Estimates of BTA Exposure Effects on Male Births



Notes: Separate event study regressions of MaleBirth on year dummies interacted with binary exposure for maternal (left), paternal (center), and provincial (right) tariff cut. Binary split at sample median. Specifications include province, birth-year, and parents' ISIC fixed effects. Vertical dashed line marks BTA implementation in late 2001. Error bars show 95% confidence intervals with province-clustered standard errors.

Figure C3: Provincial Exposure and Pre-BTA Sex Ratio at Birth



Notes: Panel (a) shows the provincial shift-share tariff exposure τ_p used to define P_p in the main specifications, displayed in quintiles. Panel (b) shows sex ratio at birth (males per 100 females) by province from 1997–1998 births in the 1999 Census. Provinces above the median exposure define the high- P_p group, and provinces above the median SRB define the high-pre-BTA-SRB subgroup in Figure 4.

Table C1: Comparison of Panel Sub-sample and Main Cross-section (2002 Baseline)

	Cross-section (2002 VHLSS)	Panel sub-sample (2002 wave)
Mother's age	28.04 (5.07)	29.05 (5.08)
Mother's education (yrs)	6.61 (3.62)	6.49 (3.60)
Ethnic minority	0.230 (0.421)	0.233 (0.423)
Urban	0.173 (0.378)	0.163 (0.369)
Province exposure	0.09 (0.02)	0.09 (0.02)
<i>N</i>	5,934	879

Notes: Standard deviations in parentheses. Cross-section uses 2002 VHLSS mothers of infants aged 0–1. Panel sub-sample uses 2002-wave observations from the rotating-panel sex-selection sample.

Table C2: Robustness to ISIC Industry Clustering

	MaleBirth (1)	log(MomWorkHrs) (2)	MomWageEmp (3)	AnyBirth (4)
$M_i \times \text{Post}$	0.08*** (0.03)	0.14*** (0.04)	0.15*** (0.03)	-0.007*** (0.002)
$F_i \times \text{Post}$	0.04 (0.02)	-0.002 (0.04)	0.02 (0.02)	0.005** (0.002)
$P_p \times \text{Post}$	0.01 (0.01)	0.06** (0.03)	-0.005 (0.01)	-0.002 (0.0009)
R^2	0.02	0.25	0.60	0.29
Observations	64,170	17,472	17,665	2,011,963
Dep. Var. Mean	0.52	4.7	0.25	0.02
Mom FE				✓
Survey Year FE		✓	✓	
Birth Year FE	✓			✓
Dad ISIC FE	✓	✓	✓	
Mom ISIC FE	✓	✓	✓	
Province FE	✓	✓	✓	
Controls	✓	✓	✓	

Notes: Dependent variables and samples follow Tables 3, 4, and 5. Controls follow the corresponding baseline table: maternal demographics, world-tariff controls, and pre-BTA sector-share-by-year controls where applicable. Standard errors are clustered by 2-digit ISIC industry (59 clusters). *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C3: Robustness to Province and ISIC Two-Way Clustering

	MaleBirth (1)	log(MomWorkHrs) (2)	MomWageEmp (3)	AnyBirth (4)
$M_i \times \text{Post}$	0.08*** (0.03)	0.14*** (0.04)	0.15*** (0.03)	-0.007*** (0.002)
$F_i \times \text{Post}$	0.04 (0.02)	-0.002 (0.03)	0.02 (0.02)	0.005** (0.002)
$P_p \times \text{Post}$	0.01 (0.01)	0.06* (0.04)	-0.005 (0.02)	-0.002 (0.001)
R^2	0.02	0.25	0.60	0.29
Observations	64,170	17,472	17,665	2,011,963
Dep. Var. Mean	0.52	4.7	0.25	0.02
Mom FE				✓
Survey Year FE		✓	✓	
Birth Year FE	✓			✓
Dad ISIC FE	✓	✓	✓	
Mom ISIC FE	✓	✓	✓	
Province FE	✓	✓	✓	
Controls	✓	✓	✓	

Notes: Dependent variables and samples follow Tables 3, 4, and 5. Controls follow the corresponding baseline table: maternal demographics, world-tariff controls, and pre-BTA sector-share-by-year controls where applicable. Standard errors are two-way clustered by province and 2-digit ISIC. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C4: Robustness to Including Potentially Endogenous Demographic Controls

	MaleBirth	log(MomWorkHrs)	MomWageEmp	AnyBirth
	(1)	(2)	(3)	(4)
$M_i \times \text{Post}$	0.08** (0.03)	0.15*** (0.05)	0.15*** (0.03)	-0.007*** (0.001)
$F_i \times \text{Post}$	0.03 (0.03)	-0.0003 (0.04)	0.02 (0.02)	0.005*** (0.001)
$P_p \times \text{Post}$	0.01 (0.02)	0.06 (0.04)	-0.004 (0.02)	-0.002 (0.002)
R^2	0.02	0.25	0.60	0.29
Observations	63,894	17,472	17,665	2,011,868
Dep. Var. Mean	0.52	4.7	0.25	0.02
Mom FE				✓
Survey Year FE		✓	✓	
Birth Year FE	✓			✓
Dad ISIC FE	✓	✓	✓	
Mom ISIC FE	✓	✓	✓	
Province FE	✓	✓	✓	
Controls	✓	✓	✓	✓

Notes: Dep. vars. are MaleBirth (col. 1), log monthly work hours (col. 2), wage employment (col. 3), and any birth (col. 4). Columns (1)–(3) add demographic controls to the corresponding main-specification control vector: mother's age and parity for MaleBirth, and mother's age for labor outcomes. Column (4) adds mother's age in the birth year; parity is not added to the fertility specification with mother fixed effects. Province-clustered standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C5: Robustness to Region-by-Year Fixed Effects

	MaleBirth	log(MomWorkHrs)	MomWageEmp	AnyBirth
	(1)	(2)	(3)	(4)
$M_i \times \text{Post}$	0.08** (0.03)	0.15*** (0.05)	0.15*** (0.03)	-0.007*** (0.001)
$F_i \times \text{Post}$	0.04 (0.03)	0.04 (0.04)	0.002 (0.02)	0.005*** (0.001)
$P_p \times \text{Post}$	0.02 (0.02)	0.05 (0.06)	-0.002 (0.02)	-0.002 (0.002)
R ²	0.02	0.26	0.61	0.29
Observations	63,894	17,472	17,665	2,011,963
Dep. Var. Mean	0.52	4.7	0.25	0.02
Mom FE				✓
Region-Year FE	✓	✓	✓	✓
Dad ISIC FE	✓	✓	✓	
Mom ISIC FE	✓	✓	✓	
Province FE	✓	✓	✓	
Controls	✓	✓	✓	✓

Notes: Binary treatment indicators (M_i , F_i , P_p) with region-by-year fixed effects added to baseline. MaleBirth and labor outcomes include province and parents' ISIC fixed effects; AnyBirth uses mother fixed effects. Controls follow the corresponding baseline table. Province-clustered standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C6: Robustness to Logit Estimation for Male Birth Outcome

	MaleBirth	
	(1)	(2)
$M_i \times \text{Post}$	0.22** (0.10)	0.31** (0.14)
$F_i \times \text{Post}$	0.11 (0.12)	0.15 (0.12)
$P_p \times \text{Post}$	0.007 (0.05)	0.05 (0.06)
Observations	64,162	64,162
Dep. Var. Mean	0.52	0.52
Birth Year FE	✓	✓
Dad ISIC FE	✓	✓
Mom ISIC FE	✓	✓
Province FE	✓	✓
Controls		✓

Notes: Logit estimates with binary treatment indicators (M_i , F_i , P_p). Specifications include birth-year, province, and parents' ISIC fixed effects. Province-clustered standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C7: Robustness to Continuous Tariff Exposure Measures

	MaleBirth				
	(1)	(2)	(3)	(4)	(5)
$\tau_m \times \text{Post}$	0.216** (0.093)		0.205** (0.090)	0.212** (0.086)	0.368** (0.142)
$\tau_d \times \text{Post}$		0.100 (0.127)	0.049 (0.126)	0.046 (0.127)	0.099 (0.131)
$\tau_p \times \text{Post}$				-0.003 (0.008)	0.002 (0.012)
R ²	0.015	0.014	0.015	0.015	0.016
Observations	66,562	66,562	66,562	66,562	64,170
Dep. Var. Mean	0.518	0.518	0.518	0.518	0.518
Controls					✓

Notes: Dep. var. is MaleBirth. Continuous treatments (τ_m , τ_d , τ_p) measure tariff reduction exposure. Columns follow the five-column progression in Table 3. Fixed effects match the baseline specification. Province-clustered standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C8: Robustness to Continuous Tariff Exposure on Maternal Labor Market Outcomes

	MomWageEmp			log(MomWorkHrs)		
	(1)	(2)	(3)	(4)	(5)	(6)
$\tau_m \times \text{Post}$	0.509*** (0.146)	0.635*** (0.153)	0.635*** (0.153)	0.146 (0.150)	0.301 (0.242)	0.326 (0.239)
$\tau_d \times \text{Post}$		-0.071 (0.074)	-0.071 (0.074)		0.148 (0.182)	0.155 (0.182)
$\tau_p \times \text{Post}$			-0.001 (0.023)			-0.086** (0.040)
R ²	0.600	0.603	0.603	0.246	0.249	0.249
Observations	17,665	17,665	17,665	17,472	17,472	17,472
Dep. Var. Mean	0.248	0.248	0.248	4.66	4.66	4.66
Controls		✓	✓		✓	✓

Notes: Dep. vars. are MomWageEmp (cols. 1–3) and log monthly hours (cols. 4–6). Continuous treatments (τ_m , τ_d , τ_p). Columns progressively add treatments and controls. Specifications include survey-year, province, and parents' ISIC fixed effects. Province-clustered standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C9: Robustness to Continuous Tariff Exposure on Fertility

	AnyBirth				
	(1)	(2)	(3)	(4)	(5)
$\tau_m \times \text{Post}$	-0.003 (0.003)		-0.008** (0.004)	-0.009*** (0.003)	-0.012** (0.005)
$\tau_d \times \text{Post}$		0.021*** (0.005)	0.023*** (0.005)	0.023*** (0.005)	0.026*** (0.005)
$\tau_p \times \text{Post}$				0.0009*** (0.0003)	0.004*** (0.001)
R ²	0.286	0.286	0.286	0.286	0.286
Observations	2,011,978	2,011,992	2,011,963	2,011,963	2,011,963
Dep. Var. Mean	0.018	0.018	0.018	0.018	0.018
Controls					✓

Notes: Dependent variable is an indicator equal to one if the mother had any birth in the past 12 months. Sample: mothers aged 20–40 from the VHLSS 2002–2008 retrospective birth panel. Continuous treatments (τ_m , τ_d , τ_p). Columns follow the same five-column progression as Table C7. All specifications include mother and birth-year fixed effects. Province-clustered standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C10: Relative>Returns Test for All Outcomes with All Four Exposure Channels

	MaleBirth	log(MomWorkHrs)	MomWageEmp	AnyBirth
	(1)	(2)	(3)	(4)
$M_i \times \text{Post}$	0.08** (0.03)	0.14*** (0.05)	0.15*** (0.03)	-0.007*** (0.001)
$F_i \times \text{Post}$	0.04 (0.03)	0.005 (0.04)	0.02 (0.02)	0.005*** (0.001)
$P_p \times \text{Post}$	0.006 (0.01)	0.11* (0.07)	-0.04 (0.03)	-0.002 (0.002)
$P_p^w \times \text{Post}$	0.007 (0.01)	-0.07 (0.07)	0.05* (0.03)	0.0008 (0.002)
R ²	0.02	0.25	0.60	0.29
Observations	64,170	17,472	17,665	2,011,963
Dep. Var. Mean	0.52	4.7	0.25	0.02
Mom FE				✓
Dad ISIC FE	✓	✓	✓	
Mom ISIC FE	✓	✓	✓	
Province FE	✓	✓	✓	

Notes: One controlled specification per outcome. M_i , F_i , P_p = binary above-median maternal, paternal, and provincial BTA exposure. P_p^w = above-median female-weighted provincial exposure ($\sum_j f_j \cdot s_{jp} \cdot \tau_j$, f_j = pre-BTA female employment share). Post = birth year > 2001 (MaleBirth, AnyBirth) or survey year > 2002 (labor). Columns (1)–(3) use the cross-sectional/labor control vector and include province, parents' ISIC, and time fixed effects. Column (4) uses the fertility panel structure from Table 5, with mother and birth-year fixed effects. Province-clustered standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C11: Sensitivity of Fertility Estimates to Retrospective Lookback Window

	AnyBirth				
	$K \leq 10$	$K \leq 12$	$K \leq 15$	$K \leq 20$	Full
	(1)	(2)	(3)	(4)	(5)
$M_i \times \text{Post}$	-0.003*	-0.005***	-0.006***	-0.007***	-0.007***
	(0.002)	(0.001)	(0.001)	(0.001)	(0.001)
$F_i \times \text{Post}$	0.006***	0.005***	0.005***	0.005***	0.005***
	(0.002)	(0.002)	(0.001)	(0.001)	(0.001)
$P_p \times \text{Post}$	-0.004*	-0.003	-0.002	-0.001	-0.001
	(0.002)	(0.002)	(0.002)	(0.002)	(0.002)
R^2	0.306	0.296	0.289	0.286	0.286
Observations	1,446,466	1,632,298	1,844,307	2,008,581	2,011,963
Dep. Var. Mean	0.018	0.018	0.018	0.018	0.018
Birth Year FE	✓	✓	✓	✓	✓
Mother FE	✓	✓	✓	✓	✓

Notes: Each column restricts the retrospective fertility panel to woman-years where the implied child age at survey is at most K years. Column (5) is the unrestricted full sample. All specifications include mother and birth-year fixed effects, world-average tariff controls, and province-clustered standard errors in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C12: Pre-Treatment Bargaining Power and Sex Selection, 1999 Census

	Mom $\leq$ Dad education	Mom $>$ Dad education
Proportion male	0.5125	0.5145
	(0.0029)	(0.0049)
N	30,548	10,619
<i>Regression. Dep. var. = Male birth</i>		
Mom $>$ Dad education		0.0026
		(0.0048)
N		41,167
Province FE		Yes

Notes: 1997–1998 births from the 1999 Vietnam Census (IPUMS). “Mom $>$ Dad education” = 1 if mother’s education exceeds father’s. Top panel reports group means (standard errors of means in parentheses). Bottom panel reports OLS with province fixed effects. Province-clustered standard errors. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

Table C13: Vietnam Infant Mortality Rate by Sex, 2002–2008

Year	Infant mortality rate			Cumulative % change	
	Male	Female	Gap	Male	Female
2002	23.1	17.4	5.7	—	—
2003	22.3	16.7	5.6	-3.5	-4.0
2004	21.6	16.2	5.4	-6.5	-6.9
2005	20.9	15.8	5.1	-9.5	-9.2
2006	20.4	15.4	5.0	-11.7	-11.5
2007	20.0	15.1	4.9	-13.4	-13.2
2008	19.6	14.8	4.8	-15.2	-14.9

Sources: World Bank, World Development Indicators. *Notes:* Infant mortality rates per 1,000 live births. “Gap” is male minus female mortality. “Cumulative % change” is relative to the 2002 baseline.